

# **Operating Systems Journal - Week 3**

## **Process Management, Scheduling, and System Performance**

---

### **Table of Contents**

1. [Process Management & System Performance](#)
  2. [Threads](#)
  3. [Concurrency Problems](#)
  4. [Problems with Race Condition Solutions](#)
  5. [Solving Race Condition](#)
  6. [Scheduling Algorithm](#)
  7. [Performance Monitoring](#)
  8. [Process Fundamentals](#)
  9. [Process Control Block \(PCB\)](#)
  10. [Process Lifecycle and State Transitions](#)
  11. [CPU Scheduling Concepts](#)
  12. [Scheduling Algorithms: FCFS and SJF](#)
  13. [Scheduling Algorithms: Priority and Round Robin](#)
  14. [Multi-threading and Process Synchronization](#)
  15. [Performance Metrics and KPIs](#)
  16. [Bottleneck Identification and Analysis](#)
  17. [Resource Utilization Patterns and Trends](#)
- 

### **1. Process Management & System Performance**

#### **Overview**

Process management is one of the fundamental responsibilities of an operating system. It encompasses the creation, scheduling, execution, and termination of processes while ensuring optimal system performance and resource utilization.

## Key Concepts

**Process Management Functions:** - **Process Creation:** The OS creates new processes through system calls (e.g., `fork()` in UNIX/Linux, `CreateProcess()` in Windows) - **Process Scheduling:** Determines which process runs on the CPU at any given time - **Process Termination:** Releases resources when a process completes or is terminated - **Inter-Process Communication (IPC):** Enables processes to communicate and synchronize

**Impact on System Performance:** 1. **CPU Utilization:** Effective process management maximizes CPU usage by minimizing idle time 2.

**Throughput:** The number of processes completed per unit time 3.

**Response Time:** Time from submission to first response 4. **Turnaround Time:** Total time from submission to completion 5. **Waiting Time:** Time spent in ready queue

## Performance Optimization Strategies

- **Context Switching Optimization:** Minimizing overhead when switching between processes
- **Load Balancing:** Distributing processes across multiple processors/cores
- **Priority Assignment:** Ensuring critical processes receive appropriate resources
- **Memory Management:** Efficient allocation and deallocation of memory resources

## Real-World Significance

Modern operating systems must manage hundreds or thousands of concurrent processes efficiently. Poor process management leads to system slowdowns, reduced responsiveness, and wasted resources.

## 2. Threads

### Overview

A thread is the smallest unit of execution within a process. While a process can contain multiple threads, all threads within a process share the same memory space and resources, making them more lightweight than separate processes.

### Thread Architecture

**Thread Components:** - **Thread ID:** Unique identifier for the thread - **Program Counter (PC):** Tracks the next instruction to execute - **Register Set:** Holds thread-specific CPU register values - **Stack:** Thread's private execution stack

**Shared Resources (among threads in same process):** - Code section - Data section - Heap memory - Open files and signals

### Types of Threads

**1. User-Level Threads (ULT)** - Managed by user-level thread libraries (e.g., POSIX Pthreads, Java threads) - Fast creation and context switching - OS is unaware of these threads - **Limitation:** If one thread blocks, entire process blocks

**2. Kernel-Level Threads (KLT)** - Managed directly by the operating system - OS can schedule threads independently - Slower context switching than ULT - **Advantage:** If one thread blocks, others can continue

**3. Hybrid Models** - Combine benefits of both ULT and KLT - Many-to-many model allows multiple user threads to map to multiple kernel threads

### Benefits of Multi-threading

**1. Responsiveness:** Application remains responsive even during lengthy operations

2. **Resource Sharing:** Threads share memory and resources, reducing overhead
3. **Economy:** Thread creation and context switching are cheaper than process operations
4. **Scalability:** Can leverage multi-core processors effectively

## Practical Example

In a web server, each client request can be handled by a separate thread, allowing the server to serve multiple clients simultaneously without creating separate processes for each connection.

---

## 3. Concurrency Problems

### Overview

Concurrency occurs when multiple processes or threads execute simultaneously or in overlapping time periods. While concurrency enables better resource utilization and performance, it introduces several critical problems that must be addressed.

### Major Concurrency Problems

1. **Race Conditions** A race condition occurs when multiple threads access shared data concurrently, and the final result depends on the timing of their execution.

**Example Scenario:** Thread 1: read balance (100) Thread 2: read balance (100) Thread 1: add 50, write balance (150) Thread 2: add 30, write balance (130) Final result: 130 (should be 180)

2. **Deadlock** A deadlock occurs when two or more processes are waiting indefinitely for resources held by each other, creating a circular wait condition.

**Four Necessary Conditions for Deadlock (Coffman Conditions): -**

**Mutual Exclusion:** Resources cannot be shared - **Hold and Wait:** Processes hold resources while waiting for others - **No Preemption:**

Resources cannot be forcibly taken from processes - **Circular Wait:**  
Circular chain of processes waiting for resources

**3. Starvation** A process is perpetually denied necessary resources because other processes are constantly given priority.

**Example:** In priority scheduling, low-priority processes may never execute if high-priority processes keep arriving.

**4. Livelock** Processes continuously change states in response to each other without making progress. Similar to deadlock but processes are actively executing, not blocked.

**5. Priority Inversion** A high-priority process is indirectly preempted by a low-priority process that holds a required resource.

## Critical Section Problem

The critical section is a code segment where shared resources are accessed. The critical section problem involves designing a protocol to ensure:

- **Mutual Exclusion:** Only one process executes in critical section at a time
- **Progress:** Selection of next process to enter critical section cannot be postponed indefinitely
- **Bounded Waiting:** Limit on how many times other processes can enter critical section before a waiting process gets its turn

## Real-World Impact

Concurrency problems can lead to data corruption, system crashes, security vulnerabilities, and unpredictable behavior. Database transactions, banking systems, and multi-user applications must carefully address these issues.

---

## 4. Problems with Race Condition Solutions

### Overview

While various solutions exist to prevent race conditions, each approach has inherent problems and limitations that must be understood for effective implementation.

## Common Solution Approaches and Their Problems

### 1. Simple Locking Mechanisms

**Problems:** - **Performance Overhead:** Acquiring and releasing locks adds computational cost - **Reduced Concurrency:** Only one thread can access protected data at a time - **Coarse-Grained Locking:** Locking large sections reduces parallelism unnecessarily

### 2. Peterson's Solution (Software-based)

**Problems:** - **Limited to Two Processes:** Not scalable to multiple processes - **Busy Waiting:** Wastes CPU cycles while waiting - **Modern Hardware Issues:** May not work correctly on modern processors due to instruction reordering - **Not Guaranteed on All Architectures:** Assumes sequential consistency

### 3. Test-and-Set / Compare-and-Swap (Hardware-based)

**Problems:** - **Busy Waiting (Spinlocks):** Continuously checks lock availability, wasting CPU - **Power Consumption:** Spinning consumes energy unnecessarily - **Cache Coherence Traffic:** Multiple cores checking same memory location causes cache traffic - **No Fairness Guarantee:** Processes may not acquire lock in order of request

### 4. Semaphores

**Problems:** - **Programming Complexity:** Easy to make errors (forget to signal, signal twice, etc.) - **No Compile-Time Checking:** Incorrect usage only detected at runtime - **Deadlock Risk:** Improper use can lead to deadlocks - **Difficult Debugging:** Race conditions are non-deterministic and hard to reproduce

### 5. Mutexes

**Problems:** - **Priority Inversion:** High-priority threads blocked by low-priority threads - **Convoy Effect:** Multiple threads waiting for one slow

**thread - Deadlock if Recursive Locking Not Supported:** Thread locks same mutex twice - **Performance Impact:** Context switching overhead when blocking

## 6. Monitors

**Problems:** - **Language Dependency:** Not all programming languages support monitors - **Hidden Complexity:** Automatic locking may hide performance issues - **Nested Monitor Problem:** Can lead to deadlocks in complex scenarios

## The Fundamental Trade-offs

**Performance vs. Correctness:** - Strict synchronization ensures correctness but reduces performance - Loose synchronization improves performance but risks race conditions

**Simplicity vs. Flexibility:** - Simple solutions (like global locks) are easy to implement but limit concurrency - Complex solutions (like fine-grained locking) improve concurrency but are error-prone

**Fairness vs. Efficiency:** - Fair scheduling prevents starvation but may reduce throughput - Efficient scheduling maximizes throughput but may cause starvation

## Best Practices

- Use the simplest synchronization mechanism that meets requirements
  - Prefer higher-level abstractions (monitors, concurrent data structures)
  - Minimize critical section size
  - Avoid holding multiple locks when possible
  - Use lock-free algorithms when appropriate
- 

## 5. Solving Race Condition

### Overview

Race conditions can be solved through various synchronization mechanisms that ensure mutual exclusion and proper ordering of operations on shared resources.

## Solution Categories

### 1. Software-Based Solutions

**Peterson's Algorithm** ``` Shared variables: boolean flag[2] = {false, false}; int turn;

Process Pi: flag[i] = true; turn = j; while (flag[j] && turn == j); // Critical Section flag[i] = false; ```

**Advantages:** - No hardware support required - Guarantees mutual exclusion, progress, and bounded waiting

**Limitations:** - Only works for two processes - Requires busy waiting - May fail on modern architectures without memory barriers

**Dekker's and Eisenberg-McGuire Algorithms:** - Earlier solutions for mutual exclusion - More complex but handle similar scenarios

### 2. Hardware-Based Solutions

**Test-and-Set Instruction** ``` boolean TestAndSet(boolean target) { boolean rv = target; \*target = true; return rv; }

while (TestAndSet(&lock)); // Critical Section lock = false; ```

**Compare-and-Swap Instruction** int CompareAndSwap(int \*value, int expected, int new\_value) { int temp = \*value; if (\*value == expected) \*value = new\_value; return temp; }

**Advantages:** - Atomic operations guaranteed by hardware - Work for any number of processes - Simple to implement

**Disadvantages:** - Busy waiting (spinlock) - Not fair - no guaranteed ordering

### 3. Semaphores

**Definition:** A semaphore is an integer variable accessed through two atomic operations: - **wait()** / **P()** / **down()**: Decrements semaphore, blocks if result is negative - **signal()** / **V()** / **up()**: Increments semaphore, may wake waiting process

**Binary Semaphore (Mutex):** ``` Semaphore mutex = 1;

```
wait(mutex); // Critical Section signal(mutex); ```
```

**Counting Semaphore:** Used for controlling access to a resource pool with multiple instances. ``` Semaphore resource\_pool = N; // N available resources

```
wait(resource_pool); // Use resource signal(resource_pool); ```
```

**Advantages:** - No busy waiting (process blocks and is added to waiting queue) - Flexible - can solve various synchronization problems - Works for any number of processes

### 4. Mutex Locks

**Definition:** A mutex (mutual exclusion lock) is a simpler version of binary semaphore specifically designed for mutual exclusion.

```
acquire_lock(); // Critical Section release_lock();
```

**Types:** - **Blocking Mutex:** Thread sleeps when lock unavailable -

**Spinlock Mutex:** Thread busy-waits when lock unavailable - **Recursive Mutex:** Same thread can lock multiple times

### 5. Monitors

**Definition:** A high-level synchronization construct that encapsulates shared data and operations with automatic mutual exclusion.

**Structure:** ``` monitor ResourceManager { // Shared data private int resource\_count;

```

// Condition variables
condition resource_available;

// Synchronized methods
public void acquire_resource() {
    while (resource_count == 0)
        wait(resource_available);
    resource_count--;
}

public void release_resource() {
    resource_count++;
    signal(resource_available);
}

} /**

```

**Advantages:** - Automatic mutual exclusion for monitor methods - Cleaner, less error-prone code - Supported by languages like Java (synchronized methods)

## 6. Lock-Free and Wait-Free Algorithms

**Concept:** Use atomic operations to modify shared data without traditional locks.

**Advantages:** - No deadlock possible - Better performance in high-contention scenarios - Immune to priority inversion

**Disadvantages:** - Complex to design and implement correctly - May use more memory (for versioning/timestamps) - Not suitable for all problems

## 7. Transactional Memory

**Concept:** Execute blocks of code atomically, similar to database transactions.

```
atomic { // Operations on shared data // Automatically
commits if no conflicts // Automatically rolls back and
retries if conflicts }
```

**Advantages:** - Composable and easier to reason about - Optimistic concurrency control - Can improve performance in low-contention

scenarios

## Selection Criteria

**Choose based on:** 1. **Complexity:** Use simplest mechanism that solves the problem 2. **Performance Requirements:** Consider overhead of different approaches 3. **Fairness Requirements:** Some mechanisms guarantee fairness, others don't 4. **Platform Support:** Hardware/language/OS support availability 5. **Scalability:** Number of concurrent threads/processes

---

# 6. Scheduling Algorithm

## Overview

CPU scheduling algorithms determine which process from the ready queue gets allocated the CPU. The choice of scheduling algorithm significantly impacts system performance, responsiveness, and fairness.

## Goals of CPU Scheduling

**Conflicting Objectives:** - **Maximize CPU Utilization:** Keep CPU as busy as possible - **Maximize Throughput:** Complete maximum number of processes per unit time - **Minimize Turnaround Time:** Reduce total time from submission to completion - **Minimize Waiting Time:** Reduce time spent in ready queue - **Minimize Response Time:** Reduce time from submission to first response - **Fairness:** Ensure all processes get fair CPU time

## Types of Scheduling

**1. Preemptive Scheduling** - CPU can be taken away from a running process - Required for time-sharing systems - Higher overhead due to context switching - Better response time

**2. Non-Preemptive Scheduling** - Process runs until completion or voluntarily yields CPU - Lower overhead - Risk of poor response time if

long processes run - Simpler to implement

## Classification of Scheduling Algorithms

Scheduling algorithms can be classified into several categories:

**Batch Systems:** - First-Come, First-Served (FCFS) - Shortest Job First (SJF) - Shortest Remaining Time First (SRTF)

**Interactive Systems:** - Round Robin (RR) - Priority Scheduling - Multi-level Queue Scheduling

**Real-Time Systems:** - Earliest Deadline First (EDF) - Rate Monotonic Scheduling (RMS) - Least Laxity First (LLF)

## Scheduling Metrics

**1. CPU Utilization** - Percentage of time CPU is actively executing processes - Target: 40-90% depending on system type

**2. Throughput** - Number of processes completed per time unit - Higher is better

**3. Turnaround Time** - Total time from process submission to completion - Turnaround Time = Completion Time - Arrival Time

**4. Waiting Time** - Total time spent in ready queue - Waiting Time = Turnaround Time - Burst Time

**5. Response Time** - Time from submission until first response - Critical for interactive systems

## Advanced Concepts

**CPU-I/O Burst Cycle** Process execution alternates between: - **CPU Burst:** Time spent executing on CPU - **I/O Burst:** Time spent waiting for I/O operations

**Dispatcher** The module that gives control of CPU to the process selected by scheduler: - Context switching - Switching to user mode -

Jumping to proper location in user program

**Dispatch Latency:** Time to stop one process and start another

## Modern Scheduling Considerations

**Multi-core Processors:** - Load balancing across cores - Processor affinity (keeping process on same core) - NUMA (Non-Uniform Memory Access) considerations

**Energy Efficiency:** - Dynamic voltage and frequency scaling (DVFS) - Scheduling to consolidate workload and allow cores to sleep

**Quality of Service (QoS):** - Different service levels for different applications - Real-time guarantees for critical processes

---

## 7. Performance Monitoring

### Overview

Performance monitoring is the continuous observation and analysis of system behavior to identify issues, optimize resource usage, and ensure system health. It provides quantitative data for making informed decisions about system tuning and capacity planning.

### Key Performance Areas

**1. CPU Performance - Metrics Monitored:** - CPU utilization percentage (overall and per-core) - Load average (1, 5, 15-minute averages) - Context switch rate - Interrupt rate - User vs. system time percentage

- **Tools:**
- top, htop (Linux)
- Task Manager, Performance Monitor (Windows)
- mpstat, vmstat, sar (detailed statistics)

**2. Memory Performance - Metrics Monitored:** - Physical memory usage (used, free, cached, buffered) - Virtual memory usage (swap) - Page fault rate - Memory allocation/deallocation rate - Cache hit/miss ratios

- **Tools:**

- `free`, `vmstat` (Linux)
- Resource Monitor (Windows)
- `pmap` (process memory maps)

**3. Disk I/O Performance - Metrics Monitored:** - Read/write throughput (MB/s) - IOPS (Input/Output Operations Per Second) - Disk utilization percentage - Average queue length - Response time/latency

- **Tools:**

- `iostat`, `iotop` (Linux)
- Performance Monitor (Windows)
- Disk Activity Monitor

**4. Network Performance - Metrics Monitored:** - Bandwidth utilization (inbound/outbound) - Packet rate - Error and drop rates - Connection count - Latency and jitter

- **Tools:**

- `netstat`, `ss`, `iftop` (Linux)
- Network Monitor (Windows)
- `tcpdump`, Wireshark (packet analysis)

**5. Process Performance - Metrics Monitored:** - Per-process CPU and memory usage - Process state (running, sleeping, zombie) - Thread count - Open file descriptors - Priority and nice values

- **Tools:**

- `ps`, `pstree` (Linux)
- Process Explorer (Windows)
- `/proc` filesystem (Linux)

## Monitoring Strategies

**1. Real-Time Monitoring** - Continuous observation of current system state - Immediate alerts when thresholds exceeded - Used for operational

troubleshooting

**2. Historical Analysis** - Collect and store metrics over time - Identify trends and patterns - Capacity planning and baseline establishment

**3. Application Performance Monitoring (APM)** - End-to-end transaction monitoring - Application-specific metrics - Code-level performance profiling

## Performance Monitoring Tools

**System-Level Tools:** - **Linux:** top, htop, atop, nmon, glances -

**Windows:** Task Manager, Resource Monitor, Performance Monitor -

**Cross-platform:** Nagios, Zabbix, Prometheus

**Application-Level Tools:** - **Profilers:** gprof, perf, valgrind (Linux), Visual Studio Profiler - **APM Platforms:** New Relic, Datadog, AppDynamics, Dynatrace

**Tracing Tools:** - strace (system call tracing) - ltrace (library call tracing) - dtrace, ftrace (kernel tracing) - BPF/eBPF (extended Berkeley Packet Filter)

## Best Practices

**1. Establish Baselines** - Understand normal system behavior - Set realistic thresholds for alerts - Document expected performance characteristics

**2. Monitor Continuously** - Don't only monitor when problems occur - Historical data helps identify root causes - Trend analysis predicts future issues

**3. Use Multiple Metrics** - Single metrics can be misleading - Correlate different metrics for complete picture - Context matters (time of day, workload type)

**4. Automate Alerting** - Set up alerts for critical thresholds - Avoid alert fatigue with proper tuning - Implement escalation policies

**5. Regular Review** - Periodic analysis of collected data - Update baselines as system evolves - Review and optimize monitoring strategy

## Common Performance Issues Detected

- **CPU Bottlenecks:** Processes competing for CPU time
- **Memory Leaks:** Gradual increase in memory usage
- **Disk I/O Bottlenecks:** Slow storage performance
- **Network Saturation:** Bandwidth limits reached
- **Resource Contention:** Multiple processes competing for resources
- **Inefficient Algorithms:** Code optimization opportunities

## Performance Tuning Cycle

1. **Monitor:** Collect performance data
  2. **Analyze:** Identify bottlenecks and issues
  3. **Tune:** Adjust configurations or code
  4. **Validate:** Verify improvements
  5. **Repeat:** Continuous improvement process
- 

## 8. Process Fundamentals

### Overview

A process is a program in execution, representing the fundamental unit of work in an operating system. Understanding process fundamentals is essential for grasping how operating systems manage and execute programs.

### Definition and Characteristics

**Process vs. Program:** - **Program:** Passive entity (executable file on disk) - **Process:** Active entity (program loaded into memory and executing) - One program can spawn multiple processes

**Process Characteristics:** 1. **Dynamic:** Created, executed, and terminated dynamically 2. **Independent:** Each process has its own

memory space

- 3. **Sequential Execution:** Instructions execute in sequence
- 4. **Controlled:** OS controls all process operations

## Process Components

### 1. Memory Layout

A process in memory consists of several sections:

```
High Memory +-----+ | Stack | <- Function calls,  
local variables | ♦♦♦ | (grows downward) +-----  
-+ | | | Free | | | +-----+ | ♦?♦ | | Heap |  
<- Dynamic memory allocation +-----+ (grows  
upward) | Data Section | <- Global and static variables |  
(initialized) | +-----+ | BSS Section | <-  
Uninitialized data +-----+ | Code/Text | <-  
Executable instructions +-----+ Low Memory
```

**Text Section:** - Contains executable code - Read-only (usually) - Shareable among processes running same program

**Data Section:** - Initialized global and static variables - Read-write permissions

**BSS (Block Started by Symbol):** - Uninitialized global and static variables - Initialized to zero at program start

**Heap:** - Dynamic memory allocation (malloc, new) - Grows upward toward higher memory addresses - Managed by programmer/runtime

**Stack:** - Function call frames - Local variables and parameters - Return addresses - Grows downward toward lower memory addresses - Managed automatically

### 2. Process Attributes

**Process Identification:** - **Process ID (PID):** Unique identifier for each process - **Parent Process ID (PPID):** PID of the process that created this process - **User ID (UID):** Owner of the process - **Group ID (GID):** Group ownership

**Process Resources:** - CPU time consumed - Memory allocation - Open files and I/O devices - Network connections - Signals and handlers

**Process Context:** - CPU register values - Program counter - Stack pointer - Memory management information - I/O status information

## Process Operations

### 1. Process Creation

**Reasons for Creating Processes:** - User initiates program execution - Batch job initialization - OS spawns process to provide service - Parent process creates child processes

**Creation Mechanisms:** - **UNIX/Linux:** `fork()` system call creates child process - Child is duplicate of parent - Different PID - Child continues execution from fork point - `exec()` family replaces child's memory with new program

- **Windows:** `CreateProcess()` creates new process
- Loads specified program into new process
- Different approach than fork/exec model

**Parent-Child Relationships:** - **Resource Sharing Options:** - Parent and child share all resources - Child shares subset of parent's resources - Parent and child share no resources

- **Execution Options:**
  - Parent and child execute concurrently
  - Parent waits until child terminates
- **Address Space:**
  - Child duplicates parent's address space
  - Child has new program loaded

### 2. Process Termination

**Normal Termination:** - Process executes last statement - Returns exit status to parent - Resources deallocated by OS

**Abnormal Termination:** - Fatal error (division by zero, invalid memory access) - Killed by another process (via signal or API) - Exceeds resource limits

**Termination System Calls:** - `exit()` (UNIX/Linux) - `ExitProcess()` (Windows)

**Parent-Child Termination:** - **Orphan Process:** Parent terminates before child - Init process (PID 1) adopts orphaned children

- **Zombie Process:** Child terminates but parent hasn't called `wait()`
- Process entry remains in process table
- Resources released but PCB remains
- Parent should call `wait()` to collect exit status

## Process Relationships

**1. Process Hierarchy** - Tree structure of processes - Init/systemd (PID 1) is ancestor of all processes in Linux - Process groups for job control

**2. Process Cooperation - Independent Processes:** Cannot affect or be affected by others - **Cooperating Processes:** Can affect or be affected by others

**Reasons for Process Cooperation:** - Information sharing - Computation speedup (parallelism) - Modularity - Convenience

**Cooperation Mechanisms:** - Shared memory - Message passing - Pipes - Sockets - Files

## Process Types

**1. CPU-Bound Processes** - Spend most time performing computations - Long CPU bursts - Examples: Scientific calculations, video encoding

**2. I/O-Bound Processes** - Spend most time performing I/O operations - Short CPU bursts followed by I/O waits - Examples: Text editors, database servers

**3. Foreground vs. Background - Foreground:** Interactive processes requiring user input - **Background:** Non-interactive processes (daemons, batch jobs)

**4. System vs. User Processes - System Processes:** OS services and daemons - **User Processes:** Application programs

## Process Advantages and Challenges

**Advantages:** - **Isolation:** Processes have separate memory spaces - **Security:** One process cannot directly corrupt another - **Stability:** Process crash doesn't affect others

**Challenges:** - **Overhead:** Process creation and context switching are expensive - **Communication:** Inter-process communication more complex than thread communication - **Resource Usage:** Each process requires separate memory allocation

---

## 9. Process Control Block (PCB)

### Overview

The Process Control Block (PCB), also known as Task Control Block (TCB), is a data structure maintained by the operating system for every process. It serves as the repository for all information needed to manage and control a process.

### Purpose and Importance

The PCB is essential for:

- 1. Process Management:** OS uses PCB to track all processes
- 2. Context Switching:** Saving and restoring process state
- 3. Scheduling:** Contains information needed by scheduler
- 4. Resource Allocation:** Tracks resources assigned to process
- 5. Inter-Process Communication:** Contains communication-related information

### PCB Structure and Contents

**1. Process Identification Information** - **Process ID (PID)**: Unique identifier for the process - **Parent Process ID (PPID)**: PID of the creating process - **User ID (UID)**: User who owns the process - **Group ID (GID)**: Group ownership

**2. Process State Information** - **Current State**: New, Ready, Running, Waiting, Terminated - **Priority**: Scheduling priority level - **Scheduling Information**: Algorithm-specific data - **CPU Burst Information**: Historical CPU usage patterns

**3. CPU Registers (Process Context)** - **Program Counter (PC)**: Address of next instruction - **Instruction Register (IR)**: Current instruction being executed - **Stack Pointer (SP)**: Top of the stack - **Base/Limit Registers**: Memory protection boundaries - **General Purpose Registers**: Data being manipulated - **Condition Code Registers**: Flags (zero, carry, overflow, etc.)

**4. Memory Management Information** - **Base and Limit Registers**: Defines memory bounds - **Page Tables**: Virtual to physical address mapping - **Segment Tables**: Segment information for segmentation - **Memory Usage Statistics**: Amount of memory allocated

**5. Process Accounting Information** - **CPU Time Used**: Total CPU time consumed - **Real Time Used**: Wall-clock time since creation - **Time Limits**: Maximum allowed CPU time - **Account Numbers**: For billing and resource tracking - **Process Start Time**: When process was created - **Time Quantum**: Remaining time slice (for RR scheduling)

**6. I/O Status Information** - **List of Open Files**: File descriptors and file pointers - **List of I/O Devices**: Allocated devices - **I/O Requests**: Pending I/O operations - **Directory Information**: Current working directory

**7. Process Privileges and Security** - **Access Rights**: Permissions for resources - **Allowed Operations**: System calls permitted - **Security Context**: SELinux context, capabilities, etc.

**8. Inter-Process Communication Information** - **Message Queues**: IPC message information - **Shared Memory Segments**: Attached shared memory - **Semaphores**: Owned or waiting semaphores - **Signals**: Pending and blocked signals

**9. Process Relationships** - **Parent Process Pointer**: Link to parent PCB  
- **Child Process Pointers**: Links to children PCBs - **Process Group ID**: For job control - **Session ID**: Terminal session identifier

## PCB Organization in OS

**Process Table**: - Array or linked list of all PCBs - Indexed by PID for quick access - OS maintains pointers to PCBs of active processes

**Queue Organization**: Different queues organize PCBs by state: - **Job Queue**: All processes in the system - **Ready Queue**: Processes ready to execute - **Device Queues**: Processes waiting for I/O devices

## PCB Operations

**1. PCB Creation** - Allocate memory for new PCB - Initialize all fields - Assign unique PID - Set initial state to "New" - Add to process table

**2. PCB Update** - Modify state during state transitions - Update CPU registers during context switch - Modify scheduling information - Update resource allocation

**3. PCB Deletion** - Called when process terminates - Free allocated resources - Remove from process table - Deallocate PCB memory

## Context Switching and PCB

### Context Switch Process:

#### 1. Save Current Process Context:

2. Save CPU register values to PCB
3. Save program counter
4. Save stack pointer

5. Update process state

#### 6. Select Next Process:

7. Scheduler selects next process from ready queue

## 8. Restore New Process Context:

9. Load CPU registers from new process PCB
10. Load program counter
11. Load stack pointer
12. Set process state to "Running"

**Context Switch Overhead:** - Time to save and restore registers - Time to update PCB - Time to switch memory maps - Cache invalidation costs

**Minimizing Context Switch Time:** - Hardware support (multiple register sets) - Efficient data structures - Minimize frequency of context switches - Use threads instead of processes where appropriate

## PCB Implementation Considerations

**Size Optimization:** - PCB size affects memory usage - Larger PCB means more context switch overhead - Balance between information stored and efficiency

**Access Speed:** - PCB frequently accessed by OS - Often kept in kernel memory - Fast lookup mechanisms (hash tables)

**Security:** - PCB contains sensitive information - Protected from user-mode access - Only kernel can modify PCB

**Portability:** - PCB structure varies across operating systems - Implementation-dependent - Abstracted through OS interfaces

## Real-World Examples

**Linux PCB (task\_struct):**

```
c struct task_struct { volatile long state; // Process state void *stack; // Kernel stack unsigned int flags; // Process flags int prio; // Priority struct mm_struct *mm; // Memory management struct files_struct *files; // Open files // ... many more fields };
```

**Windows PCB (EPROCESS):** - Contains process-specific information - Links to other structures (threads, handles, etc.) - Part of larger object manager framework

## PCB and Performance

**Impact on System Performance:** - **Memory:** PCB consumes kernel memory - **CPU:** Context switching overhead - **Scalability:** Large number of processes increases memory usage

**Optimization Strategies:** - Reduce context switch frequency - Optimize PCB layout for cache performance - Use lightweight processes (threads) when appropriate

---

## 10. Process Lifecycle and State Transitions

### Overview

The process lifecycle describes the various states a process goes through from creation to termination. Understanding state transitions is crucial for comprehending how operating systems manage process execution and resource allocation.

### Process States

- 1. New (Created)** - Process is being created - PCB is allocated and initialized - Resources are being gathered - Not yet admitted to ready queue - Transition: New  $\xrightarrow{\text{? ? ?}}$  Ready (when admitted)
- 2. Ready** - Process is prepared to execute - Waiting for CPU allocation - All necessary resources available except CPU - Resides in ready queue - Transitions: - Ready  $\xrightarrow{\text{? ? ?}}$  Running (scheduled by dispatcher) - Ready  $\xrightarrow{\text{? ? ?}}$  Suspended Ready (swapped out due to memory pressure)
- 3. Running** - Process is currently executing on CPU - Instructions being executed by processor - Only state where process actually progresses - Number of running processes  $\leq \text{number of CPU cores}$  - Transitions: - Running  $\xrightarrow{\text{? ? ?}}$  Ready (preempted or time slice expires) - Running  $\xrightarrow{\text{? ? ?}}$  Waiting (I/O request or event wait) - Running  $\xrightarrow{\text{? ? ?}}$  Terminated (completion or error)

**4. Waiting (Blocked)** - Process cannot execute until some event occurs - Waiting for I/O completion, signal, resource availability - Not consuming CPU time - Resides in appropriate waiting queue - Transitions: - Waiting  $\xrightarrow{\text{? ? ?}}$  Ready (event occurs, I/O completes) - Waiting  $\xrightarrow{\text{? ? ?}}$  Suspended Wait (swapped out)

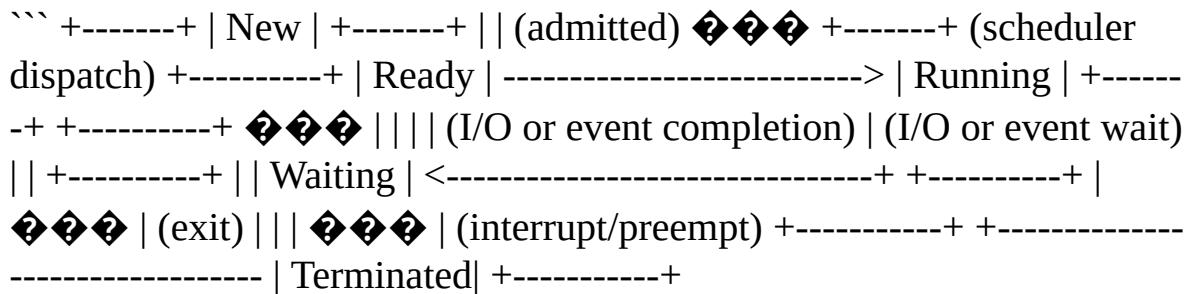
**5. Terminated (Exit)** - Process has finished execution - Exit status available for parent process - Resources being deallocated - PCB may temporarily remain (zombie state) - Final cleanup before complete removal

## Extended States (in Virtual Memory Systems)

**6. Suspended Ready** - Process is ready but swapped to disk - Frees memory for other processes - Can be resumed when memory available - Transition: Suspended Ready  $\xrightarrow{\text{? ? ?}}$  Ready (when loaded back)

**7. Suspended Wait** - Process is waiting and swapped to disk - Waiting for event while on disk - Transitions: - Suspended Wait  $\xrightarrow{\text{? ? ?}}$  Suspended Ready (event occurs) - Suspended Wait  $\xrightarrow{\text{? ? ?}}$  Waiting (loaded back while still waiting)

## State Transition Diagram



With virtual memory: Ready  $\xrightarrow{\text{? ? ?}}$  Suspended Ready Waiting  $\xrightarrow{\text{? ? ?}}$   
Suspended Wait Suspended Wait  $\xrightarrow{\text{? ? ?}}$  Suspended Ready (event  
completion) ````

## Causes of State Transitions

- 1. New ◊◊◊ Ready (Admission)** - OS admits process to execution pool - Resources available - System load permits
- 2. Ready ◊◊◊ Running (Dispatch)** - Scheduler selects process - Dispatcher performs context switch - Process allocated to CPU
- 3. Running ◊◊◊ Ready (Preemption)** - Time Slice Expiration: Quantum exhausted in round-robin - **Higher Priority Process**: Preemptive priority scheduling - **Interrupt**: System needs to handle interrupt - **Voluntary Yield**: Process calls yield()
- 4. Running ◊◊◊ Waiting (Blocking)** - I/O Request: Reading file, network operation - **Event Wait**: Waiting for signal, message - **Resource Wait**: Semaphore, mutex, resource not available - **Child Termination**: Parent waiting for child (wait() call) - **Sleep**: Process explicitly sleeps (sleep() call)
- 5. Waiting ◊◊◊ Ready (Completion)** - I/O Completion: Data available, operation finished - **Event Occurrence**: Signal received, message arrived - **Resource Available**: Lock acquired, resource freed - **Child Terminated**: wait() returns - **Timer Expiration**: Sleep time elapsed
- 6. Running ◊◊◊ Terminated (Exit)** - Normal Completion: Last instruction executed - **Error Exit**: Runtime error, exception - **Killed**: Terminated by another process or user - **Parent Termination**: Cascading termination
- 7. Ready ◊◊◊ Suspended Ready (Swapping)** - To Suspended: Memory pressure, higher priority process needs memory - **To Ready**: Memory available, process priority increased
- 8. Waiting ◊◊◊ Suspended Wait (Swapping)** - To Suspended: Free memory while waiting - **To Waiting**: Memory available, need to be responsive

## Factors Influencing Transitions

**Scheduling Policy:** - Preemptive vs. non-preemptive - Priority levels - Time quantum (for round-robin)

**System Load:** - Number of processes - Memory availability - I/O device availability

**Process Characteristics:** - CPU-bound vs. I/O-bound - Priority level - Resource requirements

**External Events:** - User actions (kill, suspend) - Hardware interrupts - System calls

## State Queues

**Ready Queue:** - Contains all ready processes - Implementation: Priority queue, FIFO queue, multi-level queue - Scheduler selects from this queue

**Device Queues:** - One queue per I/O device - Processes waiting for specific device - FIFO ordering typically used

**Job Queue:** - All processes in the system - Includes processes in all states

## Performance Implications

**State Distribution:** - **High Ready Queue:** CPU bottleneck - **High Waiting Queues:** I/O bottleneck - **Many Suspended:** Memory bottleneck

**Transition Frequency:** - Frequent state transitions increase overhead - Context switch cost impacts performance - Optimal: Balance between responsiveness and efficiency

**Queue Management:** - Efficient queue operations critical - Priority inversion prevention - Starvation avoidance

## Monitoring Process States

**Linux Commands:** bash ps aux # Shows process states (S, R, D, Z, T) top/htop # Real-time state monitoring /proc/[pid]/status # Detailed state information

**State Codes (Linux):** - **R:** Running or runnable (on run queue) - **S:** Interruptible sleep (waiting for event) - **D:** Uninterruptible sleep (usually I/O) - **T:** Stopped (by job control signal) - **Z:** Zombie (terminated but not reaped)

**Windows Task Manager:** - Running, Suspended states visible - More abstracted than Linux

## Practical Considerations

**Zombie Processes:** - Terminated but PCB remains - Parent should call wait() to reap zombie - Too many zombies can exhaust process table

**Orphan Processes:** - Parent terminated before child - Init process adopts orphans - Prevents zombie accumulation

**Process Suspension:** - User-initiated (Ctrl+Z in shell) - System-initiated (memory management) - Debugging purposes

**Long-Running States:** - Processes stuck in waiting state (D state in Linux) - May indicate hardware issues, deadlocks - Uninterruptible sleep cannot be killed

---

# 11. CPU Scheduling Concepts

## Overview

CPU scheduling is the mechanism by which the operating system determines which process in the ready queue gets access to the CPU. It is fundamental to multiprogramming and directly impacts system performance, responsiveness, and resource utilization.

## The Need for CPU Scheduling

**Multiprogramming Objective:** - Maximize CPU utilization by always having a process executing - When one process waits for I/O, another process can use the CPU - Improves overall system throughput and efficiency

**CPU-I/O Burst Cycle:** - Process execution consists of alternating CPU bursts and I/O bursts - **CPU Burst:** Period of computation - **I/O Burst:** Period waiting for I/O - Distribution of burst times affects scheduling decisions

**Burst Characteristics:** - **CPU-Bound Processes:** Long CPU bursts, infrequent I/O - **I/O-Bound Processes:** Short CPU bursts, frequent I/O - Most processes exhibit I/O-bound behavior with short CPU bursts

## Scheduling Levels

**1. Long-Term Scheduling (Job Scheduling)** - Determines which programs are admitted to the system - Controls degree of multiprogramming - Selects processes from disk/storage to load into memory - Executed infrequently (seconds, minutes) - Balances mix of CPU-bound and I/O-bound processes

**2. Medium-Term Scheduling (Swapping)** - Temporarily removes processes from memory to disk - Reduces degree of multiprogramming - Swapping decisions based on memory availability - Part of memory management

**3. Short-Term Scheduling (CPU Scheduling)** - Selects which ready process gets CPU next - Most frequent (milliseconds) - Must be extremely fast (overhead consideration) - Focus of most scheduling algorithm discussions

## Preemptive vs. Non-Preemptive Scheduling

### Non-Preemptive (Cooperative) Scheduling:

**Characteristics:** - Process keeps CPU until it terminates or blocks - No forced context switches - Simpler to implement

**When Scheduling Occurs:** 1. Process switches from running to waiting state 2. Process terminates

**Advantages:** - Lower overhead (fewer context switches) - No need for hardware timer - Simpler implementation

**Disadvantages:** - Poor response time for interactive processes - One long process can monopolize CPU - Not suitable for time-sharing systems

**Examples:** Windows 3.1, early Macintosh OS

### **Preemptive Scheduling:**

**Characteristics:** - OS can forcibly remove process from CPU - Process can be interrupted at any time - Requires hardware timer interrupt

**Additional Scheduling Points:** 3. Process switches from running to ready (preemption) 4. Process switches from waiting to ready

**Advantages:** - Better response time - Fairer CPU allocation - Essential for interactive systems

**Disadvantages:** - Higher overhead (more context switches) - Requires careful synchronization (shared data issues) - More complex implementation

**Challenges:** - **Race Conditions:** Preemption during critical section -

**Kernel Data Structures:** Protecting kernel operations - **Real-Time Systems:** Ensuring predictability

**Examples:** Modern operating systems (Linux, Windows, macOS)

## **The Dispatcher**

**Responsibilities:** 1. **Context Switch:** Save old process state, load new process state 2. **Mode Switch:** Switch from kernel mode to user mode 3.

**Jump:** Jump to proper location in user program to resume execution

**Dispatch Latency:** - Time to stop one process and start another - Should be minimized (pure overhead) - Typical range: Microseconds to

milliseconds

**Components of Dispatch Latency:** - Saving context of current process - Selecting next process (scheduling decision already made) - Loading context of selected process - Switching memory maps - Flushing and reloading cache/TLB

## Scheduling Criteria (Performance Metrics)

- 1. CPU Utilization** - Percentage of time CPU is busy - Target: 40-90% (depends on system type) - Higher utilization generally better (but not at expense of responsiveness)
- 2. Throughput** - Number of processes completed per unit time - Can be measured as processes per hour, per minute, etc. - Higher throughput indicates better performance
- 3. Turnaround Time** - Time from process submission to completion - Includes time in ready queue, executing, and waiting for I/O - Formula: Completion Time - Arrival Time - Lower is better
- 4. Waiting Time** - Total time spent in ready queue - Does not include actual execution or I/O time - Formula: Turnaround Time - Burst Time - I/O Time - Lower is better - Most commonly used metric for comparing algorithms
- 5. Response Time** - Time from submission to first response (not completion) - Critical for interactive systems - User-perceived performance metric - Formula: Time of First Response - Arrival Time - Lower is better

### Optimization Goals:

Different systems prioritize differently: - **Maximize:** CPU utilization, throughput - **Minimize:** Turnaround time, waiting time, response time

Trade-offs exist between metrics: - Minimizing average waiting time vs. minimizing variance - Throughput vs. response time - Fairness vs. performance

## Convoy Effect

**Definition:** Short processes stuck behind long process, reducing overall system performance

**Example:** - One CPU-bound process (long burst time) - Many I/O-bound processes (short burst times) - If CPU-bound process gets CPU first, others wait unnecessarily - I/O devices remain idle while CPU-bound process runs - Poor resource utilization

**Occurs in:** FCFS scheduling

**Solution:** Use preemptive scheduling algorithms (SJF, Round Robin)

## Scheduling Algorithm Comparison Framework

### Evaluation Methods:

- 1. Deterministic Modeling** - Use predefined workload - Calculate performance metrics for each algorithm - Simple but limited to specific scenarios
- 2. Queueing Models** - Mathematical analysis using queueing theory - Provides general insights - Requires assumptions about arrival and service distributions
- 3. Simulation** - Model system and processes - Run simulation with various algorithms - More realistic than deterministic modeling - Computationally intensive
- 4. Implementation** - Implement in actual system - Most accurate - Expensive and risky

## Modern Scheduling Considerations

**Multi-Core Processors:** - Load balancing across cores - Processor affinity (cache benefits) - Migration costs

**Real-Time Requirements:** - Hard real-time: Must meet deadlines - Soft real-time: Higher priority for time-critical processes

**Energy Efficiency:** - Dynamic voltage and frequency scaling (DVFS) - Core parking and sleep states - Scheduling to minimize power consumption

**Quality of Service (QoS):** - Different service levels for different applications - Resource reservations and guarantees

**Virtualization:** - Scheduling virtual machines - Two-level scheduling (hypervisor and guest OS)

---

## 12. Scheduling Algorithms: FCFS and SJF

### First-Come, First-Served (FCFS)

#### Overview

FCFS is the simplest CPU scheduling algorithm. Processes are executed in the order they arrive in the ready queue, similar to a queue at a service counter.

#### Characteristics

**Algorithm:** 1. Maintain ready queue as FIFO (First-In, First-Out) 2. When CPU becomes available, allocate it to process at head of queue 3. Process runs to completion (non-preemptive) 4. When process completes, remove from queue and select next

**Implementation:** - Simple to implement using FIFO queue - Low overhead - No starvation (every process eventually gets CPU)

**Scheduling Decisions:** - Non-preemptive: Once CPU allocated, process keeps it until completion or I/O wait - Scheduling occurs only when process terminates or blocks

#### Performance Analysis

#### Example 1: Simple Case

Processes arrive at time 0 in order P1, P2, P3:

### **Process Burst Time**

P1      24 ms

P2      3 ms

P3      3 ms

**Gantt Chart:** | P1 (0-24) | P2 (24-27) | P3 (27-30) |

**Calculations:** - Waiting Time: P1=0, P2=24, P3=27 - Average Waiting Time =  $(0+24+27)/3 = 17$  ms - Turnaround Time: P1=24, P2=27, P3=30 - Average Turnaround Time =  $(24+27+30)/3 = 27$  ms

### **Example 2: Order Matters**

Same processes arrive in order P2, P3, P1:

**Gantt Chart:** | P2 (0-3) | P3 (3-6) | P1 (6-30) |

**Calculations:** - Waiting Time: P2=0, P3=3, P1=6 - Average Waiting Time =  $(0+3+6)/3 = 3$  ms - **Significant improvement** just by changing order!

## **Advantages**

1. **Simplicity:** Easy to understand and implement
2. **Fairness:** Every process treated equally in arrival order
3. **No Starvation:** Every process eventually executes
4. **Low Overhead:** Minimal context switching

## **Disadvantages**

1. **Convoy Effect:** Short processes stuck behind long ones
2. **Poor Average Waiting Time:** Can be very high
3. **Not Optimal:** Doesn't minimize average waiting time
4. **Poor for Interactive Systems:** Long response times
5. **Order Dependent:** Performance varies greatly with arrival order
6. **No Prioritization:** Critical processes must wait like all others

## Use Cases

**Appropriate For:** - Batch systems with similar process lengths - Systems where fairness in arrival order is critical - Simple embedded systems - Background job processing

**Not Appropriate For:** - Interactive systems - Real-time systems - Time-sharing systems - Systems with mix of short and long processes

---

## Shortest Job First (SJF)

### Overview

SJF associates each process with its CPU burst length and selects the process with the smallest burst time. This algorithm is provably optimal for minimizing average waiting time.

### Characteristics

**Algorithm:** 1. Each process has associated burst time 2. When CPU available, select process with shortest burst time 3. Break ties using FCFS 4. Execute selected process

**Two Schemes:** 1. **Non-Preemptive SJF:** Once CPU given, process runs to completion 2. **Preemptive SJF (Shortest Remaining Time First - SRTF):** If new process arrives with shorter burst than remaining time of current process, preempt

### Non-Preemptive SJF

#### Example:

#### Process      Arrival Time      Burst Time

P1	0	7 ms
P2	2	4 ms
P3	4	1 ms
P4	5	4 ms

**Execution:** - At t=0: Only P1 available, execute P1 - At t=7: P2, P3, P4 in queue. P3 has shortest burst (1), execute P3 - At t=8: P2, P4 in queue. Both have burst 4, use FCFS: P2 first - At t=12: Execute P4

**Gantt Chart:** | P1 (0-7) | P3 (7-8) | P2 (8-12) | P4 (12-16) |

**Calculations:** - Waiting Time: P1=0, P2=6 (waited from 2 to 8), P3=3 (waited from 4 to 7), P4=7 (waited from 5 to 12) - Average Waiting Time =  $(0+6+3+7)/4 = 4 \text{ ms}$

## Preemptive SJF (SRTF)

### Same Example with Preemption:

**Execution:** - At t=0: P1 starts (remaining=7) - At t=2: P2 arrives (burst=4 < P1 remaining=5), preempt P1, run P2 - At t=4: P3 arrives (burst=1 < P2 remaining=2), preempt P2, run P3 - At t=5: P4 arrives (burst=4), P3 continues (remaining=0 at t=5) - At t=5: P2 resumes (remaining=2 < P4=4, P1=5) - At t=7: P2 completes, P4 runs (burst=4 < P1 remaining=5) - At t=11: P4 completes, P1 resumes - At t=16: P1 completes

**Gantt Chart:** | P1 (0-2) | P2 (2-4) | P3 (4-5) | P2 (5-7) | P4 (7-11) | P1 (11-16) |

**Calculations:** - Waiting Time: P1=9 ((11-2)+(0)), P2=1 ((2-2)+(5-4)), P3=0, P4=2 (7-5) - Average Waiting Time =  $(9+1+0+2)/4 = 3 \text{ ms}$  -

**Better than non-preemptive SJF**

## Optimality

**Proof of Optimality:** - SJF gives minimum average waiting time for a given set of processes - Moving short process before long process decreases waiting time more than it increases waiting time - Mathematical proof exists (greedy algorithm proof)

## Advantages

1. **Optimal:** Minimizes average waiting time
2. **Efficient:** Better CPU utilization than FCFS

3. **Reduced Convoy Effect:** Short processes not stuck behind long ones
4. **Better Turnaround Time:** Especially for short processes

## Disadvantages

1. **Burst Time Prediction Problem** - Cannot know exact CPU burst time in advance - Must estimate/predict future burst times - Prediction accuracy affects performance

### Prediction Methods:

**Exponential Averaging:**  $\hat{t}(n+1) = \alpha * t(n) + (1-\alpha) * \hat{t}(n)$

Where:  $\hat{t}(n+1)$  = predicted burst time for next burst  $t(n)$  = actual length of nth CPU burst  $\hat{t}(n)$  = predicted value for nth burst  $\alpha$  = weight (0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1)

- $\alpha=0$ : Recent history has no effect
- $\alpha=1$ : Only last burst matters
- Typically  $\alpha=0.5$  provides balance

**Historical Averaging:** - Use average of previous bursts - Simple but less responsive to changes

2. **Starvation** - Long processes may never execute - Continuous arrival of short processes pushes long processes back - **Solution:** Aging (gradually increase priority over time)

3. **Unfairness** - Favors short processes - Long processes suffer poor response time

4. **Implementation Complexity** - Requires burst time information - Priority queue implementation needed - Preemptive version adds context switch overhead

## Comparison: FCFS vs. SJF

<b>Aspect</b>	<b>FCFS</b>	<b>SJF</b>
Waiting Time	Can be high	Minimal (optimal)
Starvation	No	Yes (for long processes)
Complexity	Very simple	Moderate
Fairness	Fair in arrival order	Unfair to long processes
Overhead	Low	Low (non-preemptive), Medium (preemptive)
Prediction	Not needed	Required
Use Case	Batch processing	Situations where burst times known

## Practical Considerations

**Burst Time Estimation:** - Static processes (batch jobs): Use historical data - Interactive processes: Difficult to predict accurately - Mixed workload: Hybrid approaches

**Implementation in Real Systems:** - Pure SJF rarely used in general-purpose OS - Used in specialized systems (embedded, real-time) - Approximations used (priority-based with aging)

**Hybrid Approaches:** - Combine with other algorithms - Multi-level feedback queues use SJF-like behavior - Modern OS use heuristics based on process history

---

## 13. Scheduling Algorithms: Priority and Round Robin

### Priority Scheduling

#### Overview

Priority scheduling assigns a priority value to each process, and the CPU is allocated to the process with the highest priority. It's a general

framework that includes SJF as a special case (where priority is inverse of burst time).

## Characteristics

**Priority Assignment:** - **Internal Priorities:** Computed by OS based on:

- Time limits - Memory requirements - Resource needs - Ratio of CPU burst to I/O burst - Historical CPU usage

- **External Priorities:** Set by:
  - User/administrator
  - Importance of process
  - Funding or political factors
  - Department or project priority

**Priority Representation:** - Lower number = Higher priority (Unix convention) - OR Higher number = Higher priority (Windows convention) - Must be consistent within system

## Types of Priority Scheduling

**1. Non-Preemptive Priority Scheduling** - Once CPU allocated, process runs to completion - New high-priority process waits until current process finishes or blocks - Simpler implementation

**Example:**

**Process Priority Burst Time Arrival Time**

P1	3	10 ms	0
P2	1	1 ms	0
P3	4	2 ms	0
P4	5	1 ms	0
P5	2	5 ms	0

**Execution Order (lower number = higher priority):** P2 (priority 1)

◆◆◆ P5 (priority 2) ◆◆◆ P1 (priority 3) ◆◆◆ P3 (priority 4)

◆◆◆ P4 (priority 5)

**Gantt Chart:** | P2 (0-1) | P5 (1-6) | P1 (6-16) | P3 (16-18) |  
P4 (18-19) |

$$\text{Average Waiting Time} = (0+1+6+16+18)/5 = 8.2 \text{ ms}$$

**2. Preemptive Priority Scheduling** - Newly arrived high-priority process preempts running lower-priority process - More responsive to important processes - Higher context switch overhead

**Same Example with Preemption:** If processes arrive at different times, higher priority arrivals interrupt lower priority processes.

## **Major Problem: Starvation (Indefinite Blocking)**

**Issue:** - Low-priority processes may never execute - Continuous arrival of high-priority processes - **Example:** In heavily loaded system, low-priority process may wait indefinitely

**Real-World Example:** IBM System/360: Low-priority process submitted in 1973 discovered in 1980 still waiting!

## **Solution: Aging**

**Concept:** Gradually increase priority of processes waiting in system

**Implementation:** - Every X time units, increase priority by 1 - Eventually, even lowest priority process reaches highest priority - Ensures all processes eventually execute

**Example:** - Initial priority: 100 - Every 10 minutes: priority -= 1 - After 1 hour: priority = 94 - After 10 hours: reaches high priority

## **Combined Scheduling**

**Priority Classes:** Each priority level has its own queue, within which another algorithm (like Round Robin) is used

**Example:** - Priority 1 (highest): Round Robin with quantum=10ms - Priority 2: Round Robin with quantum=20ms - Priority 3 (lowest): FCFS

## **Dynamic vs. Static Priority**

**Static Priority:** - Set when process created - Never changes during execution - Simpler implementation - Risk of starvation

**Dynamic Priority:** - Changes based on behavior - Increases for I/O-bound processes (encourage responsiveness) - Decreases for CPU-bound processes (prevent monopolization) - Prevents starvation through aging - More complex but more adaptive

---

## **Round Robin (RR)**

### **Overview**

Round Robin is designed specifically for time-sharing systems. Each process gets a small unit of CPU time (time quantum/time slice), after which it's preempted and moved to the end of the ready queue.

### **Characteristics**

**Algorithm:** 1. Ready queue implemented as circular FIFO queue 2. CPU scheduler picks first process from ready queue 3. Sets timer to interrupt after 1 time quantum 4. Dispatches process 5. When quantum expires: - If process burst > quantum: Preempt, move to end of queue - If process burst  $\leq$  quantum: Process terminates or blocks, next process selected

**Key Parameter: Time Quantum (q)** - Typical range: 10-100 milliseconds - Critical for performance - Must be tuned for specific system

### **Performance Analysis**

#### **Example:**

#### **Process Burst Time**

P1      24 ms

## **Process Burst Time**

P2      3 ms  
P3      3 ms

## **Time Quantum = 4 ms**

**Gantt Chart:** | P1 (0-4) | P2 (4-7) | P3 (7-10) | P1 (10-14) |  
P1 (14-18) | P1 (18-22) | P1 (22-26) | P1 (26-30) |

**Calculations:** - P1: Turnaround =  $30-0 = 30$  ms - P2: Turnaround =  $7-0 = 7$  ms - P3: Turnaround =  $10-0 = 10$  ms - Average Turnaround Time =  $(30+7+10)/3 = 15.67$  ms

**Compare with FCFS (same processes):** FCFS Average Waiting Time = 17 ms RR Average Waiting Time = Varies but provides better response time

## **Time Quantum Selection**

**Critical Decision:** The choice of time quantum drastically affects performance

**If Time Quantum Too Large:** - Approaches FCFS behavior - Example:  $q = 1000$ ms for processes with 10ms bursts - Poor response time - Less context switching (good) but poor interactivity (bad)

**If Time Quantum Too Small:** - Excessive context switching overhead - Example:  $q = 1$ ms with 10ms context switch time - CPU spends more time switching than executing - Context switch time becomes significant portion of quantum

**Optimal Range:** - Usually 10-100 milliseconds - Context switch time should be < 10% of quantum - Typical context switch: 1-10 microseconds - Balance between responsiveness and efficiency

**Rule of Thumb:** - 80% of CPU bursts should be shorter than time quantum - Ensures most processes complete within their quantum

## **Turnaround Time Behavior**

**Counter-Intuitive Property:** Average turnaround time may be worse than SJF but response time is better

**Example showing high turnaround time:**

### **Process Burst Time**

P1	6 ms
P2	3 ms
P3	1 ms
P4	7 ms

**With q=1ms:** Excessive switching, poor turnaround time

**With q=10ms:** Each process gets full burst, essentially FCFS behavior

**With q=3ms:** Balance achieved

## **Advantages**

- 1. Fair:** Each process gets equal share of CPU
- 2. Good Response Time:** No process waits more than  $(n-1)*q$  time units ( $n$ =number of processes)
- 3. No Starvation:** Every process eventually gets CPU
- 4. Interactive:** Excellent for time-sharing systems
- 5. Simple:** Easy to implement
- 6. Predictable:** Users can estimate response time

## **Disadvantages**

- 1. Context Switch Overhead:** Quantum too small causes excessive switching
- 2. Turnaround Time:** Average turnaround time may be poor
- 3. Quantum Selection:** Difficult to choose optimal value
- 4. All Processes Equal:** Cannot prioritize important processes (unless combined with priority)
- 5. Cache Performance:** Frequent switching may reduce cache hit rates

## **Performance Variations**

**Effect of Number of Processes:** - More processes longer wait between turns - Maximum wait time:  $(n-1) * q$

**Effect of Process Behavior:** - I/O-bound processes: Often don't use full quantum (voluntarily yield) - CPU-bound processes: Always use full quantum - Mix of both: I/O-bound processes get better responsiveness

## Multi-Level Feedback Queue Extension

**Concept:** Combine RR with priorities based on execution history

**Example:** - Queue 0 (highest):  $q=8\text{ms}$  - Queue 1:  $q=16\text{ms}$  - Queue 2 (lowest): FCFS

**Behavior:** - New process starts in Queue 0 - If uses full quantum: demoted to Queue 1 - If uses full quantum again: demoted to Queue 2 - CPU-bound processes sink to lower queues - I/O-bound processes stay in high-priority queues - Adapts to process behavior automatically

## Comparison: Priority vs. Round Robin

Aspect	Priority Scheduling	Round Robin
Fairness	Unfair to low priority	Fair to all
Starvation	Possible	No
Response Time	Variable	Predictable
Implementation	Priority queue	Circular queue
Overhead	Low (non-preemptive)	Medium (preemption)
Use Case	Importance-based	Time-sharing, interactive
Flexibility	High (priority assignment)	Low (all equal)
Tuning	Priority values	Time quantum

## Real-World Usage

**Linux CFS (Completely Fair Scheduler):** - Uses red-black tree, not simple RR - Tracks virtual runtime - Each process gets proportional CPU time - Inspired by RR fairness concept

**Windows:** - Priority-based preemptive scheduling - Round-robin within each priority level - Dynamic priority adjustments

**Real-Time Systems:** - Rate Monotonic (fixed priority) - Earliest Deadline First - RR not suitable for hard real-time

## Best Practices

### 1. Time Quantum Selection:

2. Profile actual workload
3. Measure context switch time
4. Aim for 10-100ms range
5. Adjust based on system responsiveness

### 6. Hybrid Approaches:

7. Combine with priority scheduling
8. Different quantum for different priority levels
9. Adaptive quantum based on load

### 10. Monitoring:

11. Track context switch rate
  12. Monitor response time
  13. Measure CPU utilization
  14. Tune based on metrics
- 

## 14. Multi-threading and Process Synchronization

### Overview

Multi-threading enables concurrent execution within a single process, while process synchronization ensures correct cooperation between concurrent threads or processes. Together, they form the foundation for efficient and correct concurrent programming.

## **Multi-threading Fundamentals**

**Thread Benefits Revisited:** 1. **Responsiveness:** UI remains responsive during long operations 2. **Resource Sharing:** Threads share memory space naturally 3. **Economy:** Creating threads is cheaper than processes 4. **Scalability:** Utilize multiple processors effectively

**Thread Challenges:** - Shared memory requires synchronization - Race conditions possible - Debugging is more difficult - Requires careful design

## **Threading Models**

**1. Many-to-One Model** - Many user-level threads mapped to one kernel thread - Thread management done in user space - **Advantages:** Fast, no kernel involvement - **Disadvantages:** One blocking call blocks all, no true parallelism - **Example:** Green threads (older Java implementations)

**2. One-to-One Model** - Each user thread maps to kernel thread - **Advantages:** True parallelism, one blocking call doesn't block others - **Disadvantages:** Thread creation overhead, limited scalability - **Example:** Windows, Linux

**3. Many-to-Many Model** - Multiplexes many user threads onto smaller or equal number of kernel threads - **Advantages:** True parallelism, good scalability - **Disadvantages:** Complex implementation - **Example:** Some Unix variants

**4. Two-Level Model** - Combination of many-to-many and one-to-one - Allows both bound and unbound threads - **Example:** IRIX, older Solaris

## **Thread Libraries**

**POSIX Pthreads:** `c pthread_t thread; pthread_create(&thread, NULL, thread_function, arg); pthread_join(thread, NULL);`

**Java Threads:** `java class MyThread extends Thread { public void run() { // Thread code } } new MyThread().start();`

**Windows Threads:** `c CreateThread(NULL, 0, ThreadFunction, arg, 0, &threadID);`

## Process Synchronization Fundamentals

### The Critical Section Problem

**Critical Section:** Code segment where shared resources are accessed

**Requirements for Solution:** 1. **Mutual Exclusion:** Only one process in critical section at a time 2. **Progress:** If no process in critical section, selection cannot be postponed indefinitely 3. **Bounded Waiting:** Limit on times others enter critical section before waiting process gets turn

**General Structure:** `do { [Entry Section] Critical Section [Exit Section] Remainder Section } while (true);`

## Synchronization Mechanisms

### 1. Mutex Locks

**Purpose:** Simplest synchronization tool for protecting critical sections

**Operations:** ````c acquire() { while (!available) ; // busy wait available = false; }`

`release() { available = true; } ````

**Usage:** `c acquire(); // Critical Section release();`

**Characteristics:** - **Mutual Exclusion:** Guaranteed - **Busy Waiting:** Spinlock variant wastes CPU - **Simple:** Easy to understand and implement

**Variations:** - **Blocking Mutex:** Thread sleeps when lock unavailable -

**Spinlock:** Thread busy-waits (good for short critical sections) -

**Recursive Mutex:** Same thread can lock multiple times

### 2. Semaphores

**Dijkstra's Contribution:** Two types: Binary (mutex) and Counting

**Counting Semaphore:** ````c  
typedef struct { int value; struct process \*queue; } semaphore;

```
wait(semaphore *S) { S->value--; if (S->value < 0) { // Add process to S->queue block(); } }
```

```
signal(semaphore *S) { S->value++; if (S->value <= 0) { // Remove process from S->queue wakeup(process); } } ````
```

**Classic Uses:**

**Mutual Exclusion:** ````c  
semaphore mutex = 1;

```
wait(mutex); // Critical Section signal(mutex); ````
```

**Ordering (Synchronization):** ````c  
semaphore sync = 0;

```
// Process 1 statement1(); signal(sync);
```

```
// Process 2 wait(sync); statement2(); // Guaranteed to execute after statement1 ````
```

**Resource Pool:** ````c  
semaphore resource = N; // N instances of resource

```
wait(resource); // Use resource signal(resource); ````
```

### 3. Monitors

**High-Level Construct:** Combines data, operations, and synchronization

**Structure:** ````java  
monitor ResourceManager { // Shared data  
private int count = 0;

```
// Condition variables  
condition not_full;  
condition not_empty;  
  
// Synchronized operations  
public void produce() {  
    if (count == MAX)  
        wait(not_full);
```

```

        count++;
        signal(not_empty);
    }

    public void consume() {
        if (count == 0)
            wait(not_empty);
        count--;
        signal(not_full);
    }

} ***

```

**Advantages:** - Automatic mutual exclusion for monitor procedures -  
Clean abstraction - Less error-prone than explicit semaphores

**Java Implementation:**

```

```java
class BoundedBuffer {
    private int[] buffer;
    private int count = 0;

    public synchronized void insert(int item) throws InterruptedException {
        while (count == buffer.length)
            wait();
        buffer[count++] = item;
        notify();
    }

    public synchronized int remove() throws InterruptedException {
        while (count == 0)
            wait();
        notify();
        return buffer[--count];
    }
}
```

```

## Classic Synchronization Problems

### 1. Producer-Consumer (Bounded Buffer)

**Problem:** - Producers create data, place in buffer - Consumers remove data from buffer - Buffer has limited size

**Solution with Semaphores:**

```

```c
semaphore mutex = 1; // Mutual exclusion
semaphore empty = N; // Count empty slots
semaphore full = M; // Count full slots
```

```

```

0; // Count full slots

// Producer wait(empty); wait(mutex); // Add item to buffer
signal(mutex); signal(full);

// Consumer wait(full); wait(mutex); // Remove item from buffer
signal(mutex); signal(empty); ``

```

## 2. Readers-Writers

**Problem:** - Multiple readers can read simultaneously - Writers need exclusive access - No reader should wait unless writer has access

**First Readers-Writers Problem (Readers Priority):** ``c semaphore  
 mutex = 1; // Protect read\_count semaphore wrt = 1; // Writer lock int  
 read\_count = 0;

```

// Reader wait(mutex); read_count++; if (read_count == 1) wait(wrt); //  

First reader locks writers signal(mutex);

```

// Read data

```

wait(mutex); read_count--; if (read_count == 0) signal(wrt); // Last  

reader unlocks writers signal(mutex);

```

```

// Writer wait(wrt); // Write data signal(wrt); ``

```

**Problem:** Writers may starve

## 3. Dining Philosophers

**Problem:** - 5 philosophers, 5 chopsticks (forks) - Each philosopher needs 2 chopsticks to eat - Risk of deadlock if all pick up left chopstick simultaneously

**Solution 1: Odd-Even** - Odd numbered philosophers pick left first, then right - Even numbered philosophers pick right first, then left - Breaks circular wait condition

**Solution 2: At Most 4 Can Try** ``c semaphore chopstick[5] = {1, 1, 1, 1, 1}; semaphore room = 4; // At most 4 in room

```
wait(room); wait(chopstick[i]); wait(chopstick[(i+1) % 5]); // Eat  
signal(chopstick[(i+1) % 5]); signal(chopstick[i]); signal(room); ````
```

## Thread Safety

**Thread-Safe Code:** Behaves correctly when accessed from multiple threads

**Techniques:** 1. **Immutable Objects:** Cannot be modified after creation

2. **Thread-Local Storage:** Each thread has own copy

3. **Synchronization:** Use locks/semaphores/monitors

4. **Atomic Operations:** Hardware-supported atomic operations

5. **Lock-Free Algorithms:** Use compare-and-swap

**Java Example:** ````java // Not thread-safe class Counter { private int count = 0; public void increment() { count++; } // Race condition! }

```
// Thread-safe (synchronized) class Counter { private int count = 0;  
public synchronized void increment() { count++; } }
```

```
// Thread-safe (atomic) class Counter { private AtomicInteger count =  
new AtomicInteger(0); public void increment() {  
count.incrementAndGet(); } } ````
```

## Deadlock in Multi-threading

**Four Conditions (all must hold):** 1. Mutual Exclusion 2. Hold and Wait  
3. No Preemption 4. Circular Wait

**Prevention:** - Acquire locks in same order - Use timeout when acquiring locks - Detect and recover

**Best Practices:** 1. Minimize critical sections 2. Avoid nested locks

3. Use high-level abstractions (monitors, concurrent data structures)  
4. Test extensively with thread analyzers  
5. Document synchronization requirements

## Modern Synchronization Primitives

**Read-Write Locks:** - Multiple readers OR one writer - Better concurrency than simple mutex

**Condition Variables:** - Wait for specific condition - Used with mutex - Avoid busy waiting

**Barriers:** - Synchronization point for multiple threads - All threads wait until all reach barrier

**Atomic Variables:** - Lock-free operations - Compare-and-swap based - Good for counters, flags

## Performance Considerations

**Lock Contention:** - Multiple threads competing for same lock - Reduces parallelism - **Solution:** Fine-grained locking, lock-free algorithms

**Cache Coherence:** - Locks cause cache invalidations - **Solution:** Minimize shared state, use thread-local storage

**Context Switching:** - Blocking on locks causes context switches - **Solution:** Spinlocks for short critical sections

**Scalability:** - Some synchronization methods don't scale well - **Solution:** Design for scalability from start

---

# 15. Performance Metrics and KPIs

## Overview

Performance metrics and Key Performance Indicators (KPIs) provide quantitative measures of operating system and application performance. They are essential for capacity planning, troubleshooting, optimization, and ensuring Service Level Agreements (SLAs) are met.

## Categories of Performance Metrics

### 1. CPU Metrics

**CPU Utilization - Definition:** Percentage of time CPU is executing (not idle) - **Formula:**  $(\text{Total Time} - \text{Idle Time}) / \text{Total Time} \times 100\%$  - **Target Values:** - Desktop: 30-60% average - Server: 60-80% average - Critical: Investigation needed above 80% - **Interpretation:** - < 40%: Underutilized, possible waste - 40-70%: Healthy range - 70-90%: Approaching capacity - > 90%: Bottleneck likely

**Load Average - Definition:** Average number of processes in ready queue over time - **Time Periods:** 1-minute, 5-minute, 15-minute averages - **Interpretation (for single-core system):** - < 1.0: No queuing - 1.0: Fully utilized - > 1.0: Processes waiting - **Multi-core:** Multiply by number of cores - 4-core system: load of 4.0 means fully utilized

**Context Switch Rate - Definition:** Number of context switches per second - **Formula:**  $\text{Total Context Switches} / \text{Time Period}$  - **Typical Values:** 1,000 - 10,000 per second - **High Values Indicate:** - Too many processes/threads - Excessive I/O operations - Time quantum too small

**CPU Time Breakdown - User Time:** Time spent in user mode (application code) - **System Time:** Time spent in kernel mode (system calls) - **Idle Time:** Time CPU is idle - **Wait Time (I/O Wait):** Time waiting for I/O - **Steal Time (virtualization):** Time stolen by hypervisor

**Ideal Ratios:** - User: 65-70% - System: 15-25% - Idle: 10-15% - Wait: < 10%

## 2. Memory Metrics

**Memory Utilization - Definition:** Percentage of physical memory in use - **Formula:**  $(\text{Total Memory} - \text{Free Memory}) / \text{Total Memory} \times 100\%$  - **Components:** - Used: Actually used by processes - Free: Completely unused - Cached: File system cache (reclaimable) - Buffered: Kernel buffers (reclaimable)

**Available Memory** - Free + reclaimable memory - More important than "free" alone - Low available memory triggers swapping

**Page Faults - Minor Page Fault:** Page in memory but not in process page table - Fast, no disk I/O - **Major Page Fault:** Page must be loaded

from disk - Slow, significant performance impact - **Rate:** Page faults per second - **High Major Faults:** Indicates insufficient memory

**Swap Usage - Definition:** Amount of swap space used - **Swap In/Out Rate:** Pages swapped per second - **Interpretation:** - Occasional swapping: Normal - Continuous swapping (thrashing): Critical problem - High swap with low memory usage: Memory leak

**Memory Bandwidth - Definition:** Rate of data transfer to/from memory - **Measured:** GB/s - **Important for:** Memory-intensive applications

### 3. Disk I/O Metrics

**IOPS (Input/Output Operations Per Second) - Definition:** Number of read/write operations per second - **Typical Values:** - HDD (7200 RPM): 75-150 IOPS - SSD (SATA): 10,000-95,000 IOPS - NVMe SSD: 200,000-1,000,000 IOPS

**Throughput - Definition:** Amount of data read/written per second - **Units:** MB/s or GB/s - **Typical Values:** - HDD: 80-160 MB/s - SATA SSD: 200-550 MB/s - NVMe SSD: 2,000-7,000 MB/s

**Latency (Response Time) - Definition:** Time to complete single I/O operation - **Components:** - Queue time: Waiting in I/O queue - Service time: Actual I/O operation - **Typical Values:** - HDD: 10-20 ms - SSD: 0.1-1 ms - NVMe: 0.01-0.1 ms

**Disk Utilization - Definition:** Percentage of time disk is busy - **Target:** < 80% for good performance - > **90%:** Likely bottleneck

**Queue Depth - Definition:** Number of I/O requests waiting - **High Queue Depth:** Disk cannot keep up - **Optimal:** Depends on device and workload

**Read/Write Ratio - Definition:** Different workloads have different characteristics - Helps size storage appropriately

### 4. Network Metrics

**Bandwidth Utilization - Definition:** Percentage of available bandwidth in use - **Formula:** (Current Throughput / Maximum Bandwidth)  $\times 100\%$  - **Target:** < 70% average to handle bursts

**Throughput - Definition:** Actual data transfer rate - **Units:** Mbps, Gbps - **Inbound vs. Outbound:** Track separately

**Packet Rate - Definition:** Packets per second - **Important for:** Small packet workloads, firewalls

**Latency - Definition:** Time for packet to travel from source to destination - **Components:** - Propagation delay - Transmission delay - Queuing delay - Processing delay

**Packet Loss - Definition:** Percentage of packets that don't reach destination - **Acceptable:** < 1% for most applications - **Critical for:** VoIP, video streaming

**Error Rate - Definition:** Percentage of packets with errors - **Should be:** Near zero - **High errors:** Hardware problem, cable issue

**Connection Count - Definition:** Active connections - New connections per second - **Important for:** web servers, databases

## 5. Application Metrics

**Response Time - Definition:** Time from request to response - **Percentiles important:** - p50 (median): 50% of requests faster - p95: 95% of requests faster - p99: 99% of requests faster - **p99 often most important:** Captures worst user experience

**Throughput (Transactions Per Second) - Definition:** Number of transactions processed per time unit - **Units:** TPS, requests/second

**Error Rate - Definition:** Percentage of failed requests - **Target:** < 0.1% for most applications

**Availability - Definition:** Percentage of time system is operational - **Formula:** (Total Time - Downtime) / Total Time  $\times 100\%$  - **Service Levels:** - 99% ("two nines"): 3.65 days downtime/year - 99.9% ("three nines"): 8.76 hours downtime/year - 99.99% ("four nines"): 52.56

minutes downtime/year - 99.999% ("five nines"): 5.26 minutes downtime/year

**Saturation - Definition:** How "full" a resource is - Examples: Queue depths, buffer usage - Predicts future bottlenecks

## KPI Selection and Monitoring

### Choosing KPIs:

**1. Business-Critical Metrics:** - Directly impact user experience or revenue - Examples: Page load time, transaction success rate

**2. Leading Indicators:** - Predict future problems - Examples: Queue depths, cache hit rates

**3. Covering Four Golden Signals:** - **Latency:** Time to serve request - **Traffic:** Demand on system - **Errors:** Rate of failed requests - **Saturation:** How "full" system is

**SMART KPIs:** - **Specific:** Clearly defined - **Measurable:** Quantifiable - **Achievable:** Realistic targets - **Relevant:** Aligned with goals - **Time-bound:** Specific time period

## Baseline Establishment

**Process:** 1. **Collect:** Gather metrics under normal conditions 2. **Analyze:** Understand patterns (daily, weekly cycles) 3. **Document:** Record baseline values 4. **Update:** Revise as system evolves

**Benefits:** - Detect anomalies - Capacity planning - Performance regression detection - Realistic alert thresholds

## Threshold Setting

**Types:** - **Static:** Fixed value (CPU > 80%) - **Dynamic:** Based on baseline and variance - **Composite:** Multiple metrics combined

**Avoiding False Positives:** - Brief spikes normal - Use sustained thresholds (e.g., > 80% for 5 minutes) - Consider time of day (expected)

peaks)

## Performance Reporting

**Dashboard Design:** - **Real-time:** Current status at a glance - **Historical:** Trends over time - **Visual:** Charts more effective than numbers - **Layered:** Summary ♦♦♦ Detail drill-down

**Report Frequency:** - Real-time: Operations dashboard - Daily: Operations summary - Weekly: Trend analysis - Monthly: Capacity planning

## Correlation Analysis

**Why Correlate:** Single metrics can mislead; correlated metrics reveal truth

**Examples:** - High CPU + Low I/O: CPU-bound workload - High CPU + High I/O Wait: I/O bottleneck (CPU waiting) - Low CPU + High Load: I/O or memory bottleneck

**Tools:** - Statistical correlation - Time-series analysis - Multi-dimensional visualization

## SLA Metrics

**Common SLAs:** - **Availability:** 99.9% uptime - **Response Time:** 95% of requests < 200ms - **Throughput:** 10,000 TPS minimum - **Error Rate:** < 0.1% of requests

**SLA Monitoring:** - Continuous measurement - Automated alerting on violations - Regular reporting to stakeholders

## Performance Optimization Workflow

1. **Monitor:** Collect baseline metrics
2. **Identify:** Find bottleneck using metrics
3. **Hypothesize:** Form theory about cause
4. **Test:** Make targeted change

5. **Measure:** Verify improvement with metrics
  6. **Repeat:** Iterative improvement
- 

## 16. Bottleneck Identification and Analysis

### Overview

A bottleneck is a resource that limits system performance. Identifying and resolving bottlenecks is critical for system optimization. The bottleneck determines maximum system throughput; improving non-bottleneck resources has minimal impact.

### Fundamental Principles

**Theory of Constraints:** - System performance limited by its weakest link  
- Optimizing non-bottleneck resources wastes effort  
- Identify and optimize bottleneck  
- New bottleneck may emerge after resolution

**Amdahl's Law:** - Maximum speedup limited by serial portion -  
Formula: Speedup  $\leq \frac{P}{P + (1-P)N}$  - P: Proportion parallelizable  
- N: Number of processors - Implication: Even with infinite processors, speedup limited

**Little's Law:** -  $L = \lambda W$  - L: Average number in system -  $\lambda$ : Arrival rate - W: Average time in system - Helps analyze queue behavior

### Types of Bottlenecks

#### 1. CPU Bottleneck

**Symptoms:** - CPU utilization > 80-90% sustained - High load average (> number of cores) - Low I/O wait time - Processes frequently in running/ready state

**Verification:** - top, htop: Check CPU % - vmstat: Check "us" (user) and "sy" (system) columns - mpstat: Per-CPU utilization - perf top: Identify hot functions

**Common Causes:** - Inefficient algorithms ( $O(n^2)$ ) instead of  $O(n \log n)$  - Busy loops, polling - Excessive computations - Too many processes/threads competing for CPU - Single-threaded application on multi-core system

### Resolution Strategies:

**Immediate:** - Reduce process priority (nice value) - Kill unnecessary processes - Add CPU affinity to distribute load

**Short-term:** - Optimize algorithm efficiency - Parallelize workload - Enable CPU cache optimization

**Long-term:** - Add more CPU cores - Upgrade to faster processors - Distribute workload across multiple servers - Implement caching to reduce computation

## 2. Memory Bottleneck

**Symptoms:** - Available memory < 10% of total - High page fault rate (especially major faults) - Swap usage increasing - Swap in/out activity (thrashing) - Poor application performance despite low CPU usage

**Verification:** - `free -m`: Check available memory - `vmstat`: Check "si" (swap in) and "so" (swap out) - `top`: Check VIRT, RES, %MEM - `/proc/meminfo`: Detailed memory stats

**Common Causes:** - Memory leaks in applications - Too many applications running simultaneously - Poor memory management - Large data sets exceeding physical memory - Inefficient caching

### Resolution Strategies:

**Immediate:** - Kill memory-hungry processes - Clear caches (`drop_caches`) - Disable swap to prevent thrashing (if appropriate)

**Short-term:** - Fix memory leaks in code - Optimize data structures (use appropriate sizes) - Implement pagination/chunking for large datasets - Tune cache sizes

**Long-term:** - Add more physical RAM - Implement distributed memory architecture - Use memory-efficient algorithms - Consider 64-bit architecture if on 32-bit

### 3. Disk I/O Bottleneck

**Symptoms:** - High I/O wait (CPU waiting for I/O) - Disk utilization > 80% - High average queue length - High I/O latency - Slow application response despite adequate CPU/memory

**Verification:** - iostat -x: Check %util, await, avgqu-sz - iotop: Per-process I/O usage - vmstat: Check "wa" (I/O wait) - sar -d: Historical disk statistics

**Common Causes:** - Inefficient file access patterns - Too many random I/O operations - Slow storage hardware (HDD vs. SSD) - Full disk (> 90% capacity) - Fragmented filesystem - Insufficient disk cache

#### Resolution Strategies:

**Immediate:** - Identify and throttle I/O-heavy processes - Enable write-back caching (if safe) - Clear unnecessary files

**Short-term:** - Optimize I/O access patterns (sequential vs. random) - Increase readahead buffer - Tune filesystem parameters - Use faster filesystem (ext4, XFS, ZFS) - Implement I/O scheduling optimization

**Long-term:** - Upgrade to SSD or NVMe storage - Implement RAID for parallelism (RAID 0, 10) - Add dedicated I/O controllers - Separate logs, temp files, and data across drives - Use distributed storage systems

### 4. Network Bottleneck

**Symptoms:** - Network utilization > 70% sustained - High packet loss (> 1%) - Increasing latency - High retransmission rate - Application timeouts

**Verification:** - iftop: Real-time bandwidth monitoring - netstat -s: Network statistics - ss -s: Socket statistics - ping, traceroute: Latency and path analysis - tcpdump, Wireshark: Packet-level analysis

**Common Causes:** - Bandwidth saturation - Network congestion - Inefficient protocols - Excessive small packet traffic - Hardware limitations (NIC, switches) - Misconfigured network settings

### **Resolution Strategies:**

**Immediate:** - Throttle non-critical traffic - Enable QoS to prioritize critical traffic - Close unnecessary network connections

**Short-term:** - Optimize application protocols - Enable compression for data transfer - Batch small requests - Tune TCP parameters (window size, buffers) - Use CDN for static content

**Long-term:** - Upgrade network bandwidth - Use bonded/teamed network interfaces - Implement load balancing - Use faster NICs (10GbE, 40GbE) - Optimize network topology

## **5. Application Bottleneck**

**Symptoms:** - Application slow despite adequate system resources - Specific operations much slower than others - Database query slowness - Lock contention

**Verification:** - Application profiling tools - Database query analyzers - Log analysis for slow operations - Thread dump analysis

**Common Causes:** - Inefficient database queries (missing indexes, full table scans) - N+1 query problem - Lock contention, deadlocks - Poor algorithm choice - Synchronous operations that could be asynchronous - External API slowness

### **Resolution Strategies:**

**Immediate:** - Kill slow queries - Add missing database indexes - Increase connection pool size

**Short-term:** - Optimize SQL queries - Implement caching (Redis, Memcached) - Use connection pooling - Reduce lock scope and duration - Implement asynchronous processing

**Long-term:** - Refactor application architecture - Implement microservices - Use message queues for async operations - Database sharding/partitioning - Implement read replicas

## Bottleneck Identification Methodology

### Systematic Approach

- 1. Establish Baseline** - Normal operation metrics - Expected behavior patterns - Historical trends
- 2. Define Problem** - What is slow? - When does slowness occur? - Who is affected? - Quantify: How slow? (response time, throughput)
- 3. Gather Data** - System metrics (CPU, memory, disk, network) - Application metrics - Logs and traces - User reports
- 4. Analyze Patterns** - Time-based patterns (time of day, day of week) - Load correlation (slowness vs. user count) - Resource correlation (which resources saturated)
- 5. Form Hypothesis** - Based on data analysis - Identify most likely bottleneck - Consider dependencies
- 6. Test Hypothesis** - Focused monitoring - Controlled tests - Reproduce issue in test environment if possible
- 7. Implement Solution** - Start with lowest-cost, highest-impact - Change one thing at a time - Document changes
- 8. Verify Resolution** - Measure performance improvement - Monitor for side effects - Ensure problem doesn't recur
- 9. Iterate** - New bottleneck may emerge - Continuous improvement

### Tools and Techniques

#### System-Level Tools:

**Linux:** ` ` ` bash

# CPU

top, htop, mpstat, perf

# Memory

free, vmstat, pmap

# Disk I/O

iostat, iotop, fio

# Network

iftop, nethogs, ss, tcpdump

# Comprehensive

sar (system activity reporter) atop (all-in-one) ```

**Windows:** - Performance Monitor (perfmon) - Resource Monitor - Task Manager - Process Explorer (Sysinternals)

**Profiling Tools:** - perf (Linux kernel profiling) - valgrind (memory profiling) - gprof (code profiling) - strace/ltrace (system/library call tracing)

**APM (Application Performance Monitoring):** - New Relic - Datadog - AppDynamics - Dynatrace

**Specialized:** - Database: Explain plans, slow query logs - Web: Browser dev tools, Lighthouse - Network: Wireshark, tcpdump

## Common Patterns

**Pattern 1: CPU Bound** - High CPU, low I/O wait - Action: Optimize code, add cores, distribute load

**Pattern 2: I/O Wait** - Low CPU user time, high I/O wait - Check: Disk stats, network stats - Action: Optimize I/O, upgrade storage/network

**Pattern 3: Memory Pressure** - High memory usage, swapping - Action: Add RAM, fix leaks, optimize usage

**Pattern 4: Lock Contention** - Low CPU despite many threads, threads in waiting state - Action: Reduce lock scope, use finer-grained locking, lock-free algorithms

**Pattern 5: External Dependency** - System resources fine, but slow responses - Action: Optimize external calls, implement caching, use async

## Performance Testing for Bottleneck Discovery

**Load Testing:** - Gradually increase load - Identify breaking point - Find which resource saturates first

**Stress Testing:** - Push system beyond normal capacity - Identify failure modes - Determine absolute limits

**Spike Testing:** - Sudden load increase - Test auto-scaling, caching effectiveness

**Endurance Testing:** - Sustained load over extended period - Identify memory leaks, resource exhaustion

## Documentation and Communication

**Bottleneck Report Should Include:** 1. **Problem Description:** What is slow? 2. **Impact:** Who is affected? Business impact? 3. **Evidence:** Metrics showing bottleneck 4. **Root Cause:** Why is this a bottleneck? 5. **Solution:** Recommended fix 6. **Cost/Effort:** Resources required 7. **Expected Improvement:** Quantified benefit

**Prioritization:** - Impact (users affected, business value) - Effort (implementation complexity, cost) - Risk (stability, side effects)

## Prevention

**Proactive Measures:** 1. **Capacity Planning:** Predict future needs 2.

**Regular Monitoring:** Detect issues early 3. **Performance Testing:** Before production deployment 4. **Code Reviews:** Focus on performance 5. **Architecture Reviews:** Identify design bottlenecks 6. **Alerting:** Automated detection of saturation 7. **Regular Optimization:** Don't wait for crisis

---

# 17. Resource Utilization Patterns and Trends

## Overview

Understanding resource utilization patterns and trends enables proactive capacity planning, predictive maintenance, and efficient resource allocation. Analyzing trends over time reveals insights that point-in-time snapshots cannot provide.

## Temporal Patterns

### 1. Time-of-Day Patterns

**Typical Business Application:** - **Morning Spike** (8-10 AM): Users logging in, batch jobs completing - **Lunch Dip** (12-1 PM): Reduced activity - **Afternoon Peak** (2-4 PM): Maximum concurrent users - **Evening Decline** (6 PM+): Cleanup jobs, backups

**Implications:** - Size resources for peak, not average - Schedule maintenance during low-usage periods - Plan batch jobs around user activity

**Analysis Techniques:** - Hour-of-day aggregation - Identify peak hours - Calculate peak-to-average ratio

## 2. Day-of-Week Patterns

**Business Systems:** - **Monday:** Often highest load (weekend backlog) - **Weekdays:** Consistent high usage - **Friday Afternoon:** Declining usage - **Weekend:** Low usage (opportunity for maintenance)

**Consumer Systems:** - **Weekends:** Peak usage for entertainment, shopping - **Evenings:** Higher than daytime

**Analysis:** - Compare weekday vs. weekend patterns - Plan deployments and maintenance - Identify anomalies (unusual weekend spikes)

## 3. Seasonal Patterns

**Retail:** - Holiday shopping seasons (Black Friday, Christmas) - Back-to-school periods - Tax season for accounting software

**Educational:** - Registration periods - Exam schedules - Summer drop-off

**Business Cycles:** - End-of-quarter reporting - Annual planning cycles - Fiscal year-end

**Analysis:** - Year-over-year comparisons - Seasonal decomposition - Capacity planning for known events

## 4. Growth Trends

**Linear Growth:** - Steady user acquisition - Predictable resource needs - Formula:  $y = mx + b$

**Exponential Growth:** - Viral adoption - Network effects - Formula:  $y = a \diamond \diamond e^{bx}$  - Requires aggressive scaling

**Logarithmic Growth:** - Market saturation - Mature products - Growth slowing over time

**Analysis:** - Trend line fitting - Projection into future - Determine when capacity limits reached

## Resource-Specific Patterns

### CPU Utilization Patterns

**Burst Pattern:** - Short spikes of high CPU - Long periods of low usage - **Characteristics:** Web servers handling requests - **Optimization:** Ensure quick response to bursts

**Sustained Pattern:** - Consistent high CPU usage - **Characteristics:** Batch processing, scientific computation - **Optimization:** Maximize throughput

**Periodic Pattern:** - Regular spikes at predictable intervals - **Characteristics:** Scheduled jobs, cron tasks - **Optimization:** Spread jobs to avoid simultaneous execution

**Trending Up:** - Gradually increasing baseline - **Cause:** Growing user base, inefficient code (memory leak causing GC) - **Action:** Investigate root cause, plan capacity increase

### Memory Utilization Patterns

**Steady State:** - Relatively constant after startup - **Healthy:** Indicates stable application - **Monitoring:** Watch for gradual increases

**Sawtooth Pattern:** - Memory increases, then drops (garbage collection) - **Normal:** For garbage-collected languages (Java, Python) - **Problem:** If peaks increasing or valleys not returning to baseline (memory leak)

**Staircase Pattern:** - Step increases in memory usage - **Cause:** New processes/services starting, data loading - **Monitoring:** Ensure stairs don't exceed capacity

**Continuous Increase:** - Memory usage never decreases - **Problem:** Memory leak - **Action:** Profile application, fix leak

### Disk I/O Patterns

**Random Access Pattern:** - Scattered reads/writes across disk - **Characteristics:** Database OLTP workloads - **Optimization:** Use SSD,

increase cache

**Sequential Pattern:** - Contiguous reads/writes - **Characteristics:** Log files, video streaming, backups - **Optimization:** HDD acceptable, optimize buffer size

**Write-Heavy Pattern:** - More writes than reads - **Characteristics:** Logging systems, data collection - **Optimization:** Write-optimized storage (LSM trees), SSD with good write endurance

**Read-Heavy Pattern:** - More reads than writes - **Characteristics:** Web content delivery, analytics - **Optimization:** Caching, read replicas

## Network Traffic Patterns

**Bursty Traffic:** - Short periods of high bandwidth - **Characteristics:** File downloads, backups - **Optimization:** Ensure bandwidth headroom for bursts

**Streaming Traffic:** - Sustained moderate bandwidth - **Characteristics:** Video streaming, VoIP - **Optimization:** QoS, latency minimization

**Microbursts:** - Very short (milliseconds) high-bandwidth spikes - **Problem:** Can cause packet loss even with low average utilization - **Detection:** Requires high-resolution monitoring

**Asymmetric Traffic:** - Inbound >> Outbound (download-heavy) - Outbound >> Inbound (upload-heavy, backup to cloud) - **Planning:** Size bandwidth appropriately for direction

## Correlation Patterns

### User Activity Correlation

**Direct Correlation:** - Resource usage proportional to user count - **Example:** Web server CPU vs. concurrent users - **Use:** Capacity planning based on user growth

**Threshold Behavior:** - Linear until threshold, then exponential - **Example:** Database performance vs. connection count - **Use:** Identify

and avoid threshold

**Inverse Correlation:** - One resource up, another down - **Example:** Cache hit rate up, disk I/O down - **Use:** Validate optimization effectiveness

## Multi-Resource Patterns

**CPU-Memory Trade-off:** - Caching: More memory less CPU (recomputation) - Compression: More CPU less memory/disk/network

**I/O-CPU Trade-off:** - Caching: More memory less I/O, more CPU (cache management)

**Balanced Utilization:** - All resources ~70% utilized - Indicates well-architected system

**Single Resource Saturated:** - One resource at 100%, others low - Indicates bottleneck (see section 16)

## Anomaly Detection

### Types of Anomalies

**Point Anomaly:** - Single data point significantly different - **Example:** CPU spike for 1 second - **Cause:** Often transient, may ignore if rare

**Contextual Anomaly:** - Normal in one context, anomalous in another - **Example:** High traffic normal at noon, anomalous at 3 AM - **Detection:** Requires time-of-day context

**Collective Anomaly:** - Individual points normal, but sequence is not - **Example:** Gradual CPU increase over hours - **Detection:** Trend analysis

## Anomaly Detection Techniques

**Statistical Methods:** - **Standard Deviation:** Values  $> 3\sigma$  from mean - **Percentile-Based:** Values outside 1st-99th percentile - **Moving**

**Average:** Compare to rolling average

**Machine Learning:** - **Clustering:** Identify outliers from clusters -  
**Isolation Forest:** Isolate anomalies - **LSTM/RNN:** Learn temporal patterns, detect deviations

**Threshold-Based:** - **Static Threshold:** Fixed value (CPU > 80%) -  
**Dynamic Threshold:** Adjust based on historical data

**Rate of Change:** - Detect sudden changes - **Example:** Memory increase > 10% in 5 minutes

## Trend Analysis Techniques

### Moving Averages

**Simple Moving Average (SMA):** - Average over fixed window - Smooths short-term fluctuations - Lags current value

**Exponential Moving Average (EMA):** - Weighted average, more weight to recent values - More responsive than SMA - Formula:  $EMA_t = \alpha \cdot Value_t + (1-\alpha) \cdot EMA_{t-1}$

**Use Cases:** - Smooth noisy data - Identify underlying trends - Alerting (compare current to moving average)

### Seasonal Decomposition

**Components:** 1. **Trend:** Long-term direction 2. **Seasonal:** Regular periodic variation 3. **Residual:** Random noise

**Methods:** - **Additive:**  $Value = Trend + Seasonal + Residual$  -  
**Multiplicative:**  $Value = Trend \cdot Seasonal \cdot Residual$

**Use Cases:** - Remove seasonality to see true growth - Forecast future values - Identify abnormal patterns

### Regression Analysis

**Linear Regression:** - Fit line to data - Predict future values - R<sup>2</sup>? value indicates fit quality

**Polynomial Regression:** - Fit curve to data - Better for non-linear trends

**Multiple Regression:** - Multiple variables predict one variable -  
**Example:** CPU = f(users, requests, data\_size)

## Forecasting

### Time Series Forecasting

**Simple Methods:** - **Last Value:** Use most recent value - **Average:** Use historical average - **Trend Projection:** Extend linear trend

**Advanced Methods:** - **ARIMA:** Autoregressive Integrated Moving Average - **Prophet:** Facebook's forecasting tool, handles seasonality - **LSTM:** Deep learning for complex patterns

**Forecast Horizon:** - Short-term (minutes-hours): High accuracy possible - Medium-term (days-weeks): Moderate accuracy - Long-term (months-years): Lower accuracy, trends only

## Capacity Planning

### Forecasting Resource Needs:

1. **Collect Historical Data:** At least 3-6 months
2. **Identify Trends:** Growth rate, seasonality
3. **Project Forward:** Using appropriate model
4. **Add Safety Margin:** 20-30% buffer
5. **Plan Acquisition:** Lead time for procurement
6. **Monitor and Adjust:** Revise forecast regularly

**Example:** - Current CPU usage: 60% - Growth rate: 5% per month - Months to 80% (upgrade threshold): 4 months - Procurement lead time: 2 months - **Action:** Order upgrades now

## Visualization Techniques

**Time Series Graphs:** - Line charts for continuous metrics - Multiple series for comparison - Annotations for events (deployments, outages)

**Heatmaps:** - Time of day vs. day of week - Visualize patterns at a glance - Color intensity indicates value

**Histograms:** - Distribution of values - Identify normal range vs. outliers

**Scatter Plots:** - Correlation between two metrics - Identify relationships

**Dashboards:** - Combine multiple visualizations - Real-time and historical views - Drill-down capability

## Best Practices

**Data Collection:** - **Granularity:** Balance detail vs. storage - Real-time: 1-second intervals - Short-term (days): 1-minute intervals - Long-term (months): 5-minute or hourly averages - **Retention:** - High-resolution: 7-30 days - Aggregated: 1-2 years - **Consistency:** Regular sampling intervals

**Analysis:** - **Context:** Always consider time context (day/week/season) - **Multiple Metrics:** Don't rely on single metric - **Baseline Comparison:** Compare to established baseline - **Correlation:** Look for relationships between metrics

**Alerting:** - **Avoid False Positives:** Use appropriate thresholds - **Context-Aware:** Different thresholds for different times - **Sustained Issues:** Alert on sustained, not momentary spikes - **Rate of Change:** Alert on sudden changes, not just levels

**Documentation:** - **Pattern Library:** Document known patterns - **Event Log:** Correlate with deployments, incidents - **Lessons Learned:** Document anomalies and causes

**Automation:** - **Automated Reporting:** Regular trend reports - **Anomaly Detection:** Automated alerts - **Capacity Forecasting:** Automated projections - **Auto-Scaling:** Respond to patterns automatically

## Case Study Example

**Scenario:** Web application experiencing slow response times

**Analysis:** 1. **Collect Data:** CPU, memory, disk, network over past month  
2. **Identify Pattern:** CPU spikes every day at 2 PM  
3. **Correlate:** 2 PM coincides with automated report generation  
4. **Root Cause:** Report queries inefficient, full table scans  
5. **Solution:** Add database indexes, optimize queries  
6. **Validate:** Monitor CPU pattern after fix  
7. **Result:** 2 PM spikes eliminated, average CPU reduced 20%

**Trend Observation:** Overall CPU baseline increased 15% over 3 months

**Forecast:** At current growth rate, will reach 80% utilization in 6 months

**Action:** - Immediate: Optimize application (reduces baseline) - Short-term: Plan database optimization (vertical scaling) - Long-term: Plan horizontal scaling architecture

---

## Conclusion

This comprehensive journal covers the essential topics in operating systems process management, scheduling, concurrency, and performance monitoring. Understanding these concepts is crucial for:

1. **System Administration:** Efficiently managing and troubleshooting operating systems
2. **Software Development:** Writing efficient, concurrent applications
3. **Performance Engineering:** Optimizing system and application performance
4. **Capacity Planning:** Predicting and preparing for future resource needs

Each topic builds upon previous ones, creating a cohesive understanding of how modern operating systems manage processes, schedule tasks, handle concurrency, and maintain optimal performance.

## Key Takeaways

- **Process management** is fundamental to operating system functionality
- **Scheduling algorithms** must balance competing objectives (throughput, response time, fairness)
- **Concurrency** enables better resource utilization but requires careful synchronization
- **Performance monitoring** provides visibility into system behavior and health
- **Bottleneck identification** focuses optimization efforts where they matter most
- **Trend analysis** enables proactive management and capacity planning

## Further Study

To deepen your understanding:

- Implement scheduling algorithms in code
- Practice using performance monitoring tools
- Analyze real-world system performance issues
- Study advanced topics (real-time scheduling, multi-core scheduling, virtualization)

---

**Journal Prepared for:** Week 3 Operating Systems Coursework  
**Topics Covered:** 17 comprehensive sections on Process Management and System Performance  
**Academic Level:** University-level detailed analysis  
**Focus:** Higher marks achievement through comprehensive understanding and detailed explanations

# Table of Contents

|   |    |
|---|----|
| Process Management & System Performance         | 1  |
| Threads   | 3  |
| Concurrency Problems                            | 4  |
| Problems with Race Condition Solutions          | 5  |
| Solving Race Condition                          | 7  |
| Scheduling Algorithm                            | 11 |
| Performance Monitoring                          | 13 |
| Process Fundamentals                            | 16 |
| Process Control Block (PCB)                     | 20 |
| Process Lifecycle and State Transitions         | 24 |
| CPU Scheduling Concepts                         | 28 |
| Scheduling Algorithms: FCFS and SJF             | 33 |
| Scheduling Algorithms: Priority and Round Robin | 38 |
| Multi-threading and Process Synchronization     | 45 |
| Performance Metrics and KPIs                    | 52 |
| Bottleneck Identification and Analysis          | 58 |
| Resource Utilization Patterns and Trends        | 65 |