

Inter-Process Communication (IPC) Mechanisms

Operating Systems Lecture



Processes



Communication



Resources

Learning Focus

L02: IPC mechanisms and theoretical foundations

Instructor: Tanaya Bowade & Dr Shabih Fatima

Course: Operating Systems

Introduction to IPC

i Definition

Inter-Process Communication (IPC) is a fundamental mechanism within an operating system that enables COOPERATING processes to exchange data and coordinate their activities.

?

Why Processes Need Communication

Information Sharing

Multiple processes accessing shared data

Modularity

Enabling modular components to interact

Computation Speedup

Breaking tasks into concurrent subtasks

Convenience

Structured way for processes to interact

Real-world Examples



Web Browsers

Separate processes for tabs/extensions communicating to share data



Shell Pipelines

Commands chained together using pipes (e.g., ls | grep .txt)



Client-Server Apps

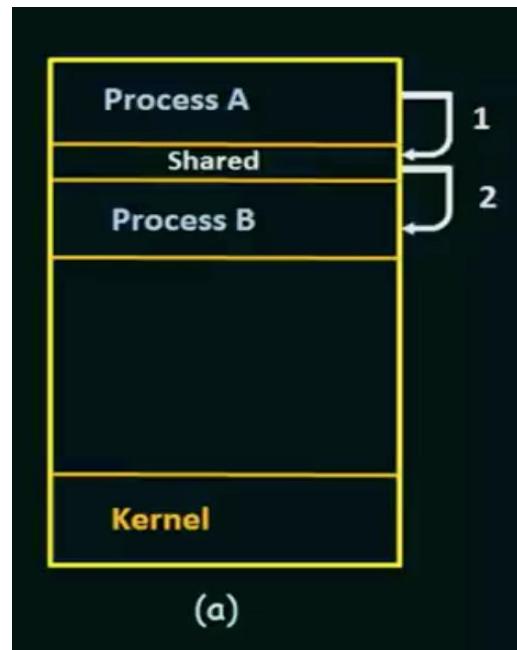
Client sends requests to server, which processes and responds

IPC Models Overview

Inter-Process Communication mechanisms rely on two primary models: **shared memory** and **message passing**. The choice between them depends on application requirements.

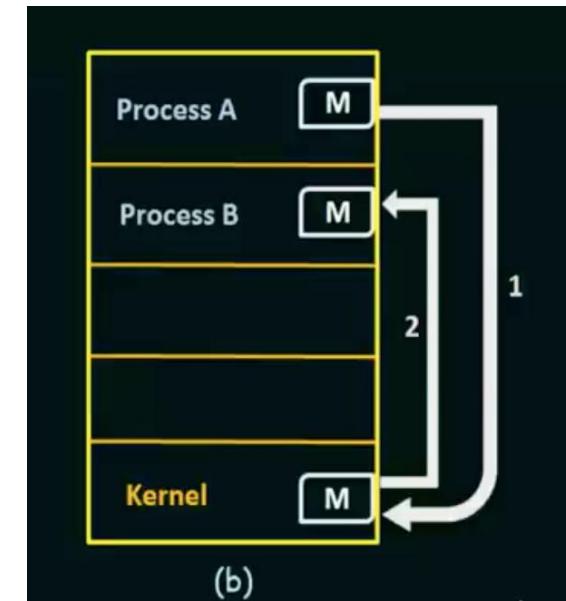
Shared Memory Model

- Speed**
High speed - direct memory access
- Complexity**
Higher - requires synchronization
- Synchronization**
Explicit needed for data consistency



Message Passing Model

- Speed**
Lower - kernel involvement
- Complexity**
Lower - simpler to manage
- Synchronization**
Built-in to message semantics

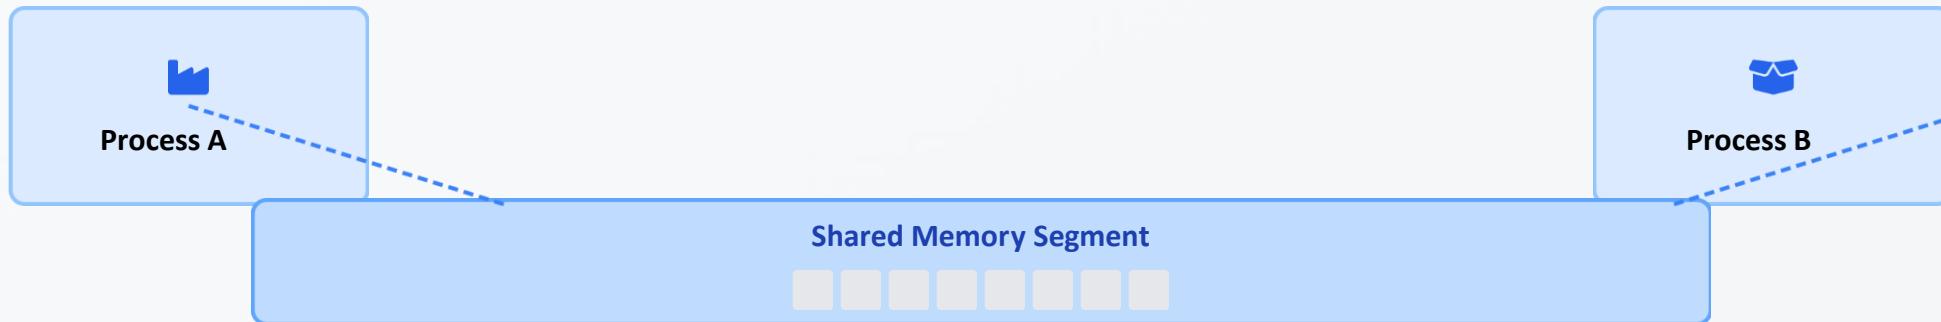


Key Decision Factor: Choose shared memory for performance, message passing for simplicity.

Shared Memory Systems

Concept

Shared memory allows multiple processes to access the same physical memory region for high-speed data exchange.



+ Advantages

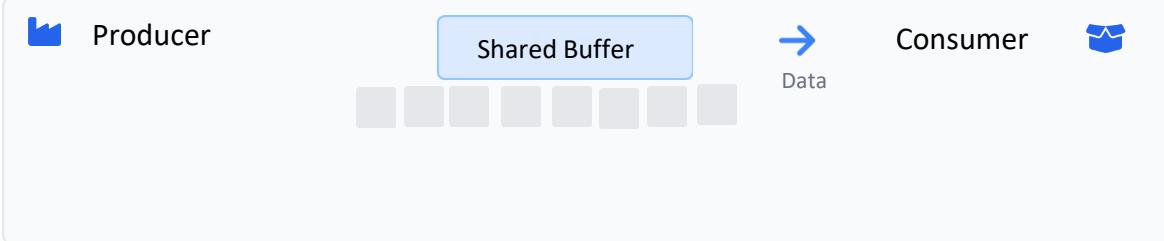
- ⚡ **High Speed**
Direct memory access without system calls
- ↔ **Efficient for Large Data**
No data copying between address spaces

⚠ Challenges

- ⟳ **Complex Synchronization**
Requires semaphores, mutexes to prevent race conditions
- 🛡 **Security Risks**
Data vulnerable if not properly managed

⚙ Producer-Consumer Example

- A classic example of shared memory where:
- Producer generates data into shared buffer
 - Consumer retrieves data from same buffer
 - Synchronization ensures producer doesn't write to full buffer
 - Consumer doesn't read from empty buffer



Message Passing Systems

Concept

Message passing is an IPC mechanism where processes communicate by exchanging messages. Communication occurs through the operating system kernel via system calls such as `send()` and `receive()`.

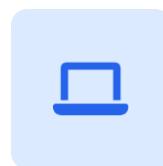
→ Direct Communication

- ✓ Processes explicitly name each other
- </> `send(P, message)`: sends to process P
- </> `receive(Q, message)`: receives from process Q

✉️ Indirect Communication

- ✓ Uses mailboxes/ports to facilitate communication
- ✓ Messages sent to specific mailboxes
- ✓ Mailboxes can be shared by multiple processes

Client-Server Example



Client Process
(Browser)

Request Message
`send(server, request)`

Response Message
`send(client, response)`



Server Process
(Web Server)

A web browser (client) sends a request message to a web server. The server processes the request and sends a response back. This exchange relies on message passing mechanisms.

Pipes and Redirection in Unix

💡 Unix Philosophy

Unix emphasizes building small, simple, and modular programs that each do one thing well, combining them using mechanisms like pipes for complex tasks.

▶ Pipe (|)

Connects stdout of one command to stdin of another:

```
ls -l | wc -l
```

Lists files then counts lines (files/directories)

▶ Redirection

Output (>):

```
echo "Hello" > file.txt
```

Writes "Hello" to file.txt (overwrites if exists)

Input (<):

```
sort < names.txt
```

Sorts lines in names.txt

▶ How Pipes Work

>_ Cmd1



stdout
Produces output

Pipe

>_ Cmd2



stdin
Consumes input

Pipes, FIFOs, and Sockets

↔ Pipes vs. FIFOs

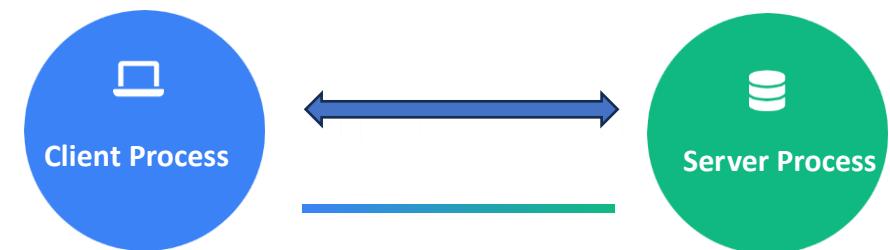
Feature	Pipes (Anonymous)	FIFOs (Named)
Nature	Unidirectional communication	One-way (half-duplex) — but can be opened twice for full-duplex communication
Naming	Unnamed; created in memory	Named; appears as a special file in the file system
Relationship	Used between related processes (e.g., parent-child)	Used between unrelated processes
Persistence	Exist only as long as communicating processes are active	Persist until explicitly deleted from the file system
Creation	Created using the `pipe()` system call	Created using the `mkfifo()` system call

🌐 Sockets

Sockets serve as an endpoint for communication, enabling processes to exchange data either on the same machine or across a network.

- ✓ More versatile and robust IPC mechanism compared to pipes and FIFOs
- ✓ Support various communication protocols (TCP, UDP)
- ✓ Unix domain sockets allow IPC between processes on the same machine

Socket Communication Example



Client-server model: Fully bidirectional communication channel

Signals and Process Communication

Definition

Signals are asynchronous notifications in Unix-like systems that indicate events to processes.

Common Signals

SIGINT

Generated by Ctrl+C. Requests graceful termination.

SIGKILL

Forceful termination. Cannot be caught or ignored.

SIGTERM

Request for termination. Can be caught and handled.

How Processes Use Signals

Process Management

- Parent sends signals to children
- Child notifies parent of events

System Calls

`kill()`

Synchronization Primitives

🔒 Introduction

Synchronization primitives are crucial for managing concurrent access to shared resources in multi-process or multi-threaded environments. They ensure data consistency and prevent race conditions.

Semaphores

- ✓ Signaling mechanisms for controlling access to resources
- ✓ Maintain a count of available resources
- ✓ Process can block until resource is available



Mutexes

- ✓ Mutual exclusion objects for exclusive access
- ✓ Ensures only one process/thread accesses critical section
- ✓ Process blocks until mutex is available



Monitors

- ✓ High-level constructs encapsulating shared data
- ✓ Provide mutual exclusion for data access
- ✓ Support condition synchronization

👉 Importance in IPC

- 🛡 Prevent race conditions in shared data
- ⚖ Ensure data consistency across processes
- ⌚ Coordinate process execution order



IPC in Practice

Real-world Applications



Databases

Utilize shared memory for efficient cache management, allowing multiple database processes to access and modify shared data structures in memory, reducing disk I/O and improving performance.



Web Servers

Employ multiple processes or threads to handle concurrent client requests. These processes communicate using various IPC mechanisms, such as message queues or sockets, to coordinate tasks and share session information.



OS Components

The kernel and various user-space daemons communicate extensively using IPC. For instance, a system logger might receive messages from different processes via message queues or named pipes.



Microservices

Distributed systems often rely on IPC, particularly network sockets (TCP/UDP), for communication between independent services running on the same or different machines.

Performance Considerations



Kernel Involvement

Message passing has higher overhead due to context switches and data copying, while shared memory allows direct memory access.



Data Amount

For large data transfers, shared memory is more efficient as it avoids copying data between process address spaces.



Process Location

Local communication favors shared memory, while networked communication typically requires message passing mechanisms.



Communication Frequency

For frequent communication, the overhead of establishing connections in message passing may be offset by its simplicity.

Review: IPC Theoretical Foundations

Key Concepts Summary

Shared Memory

Direct memory access for high-speed data exchange, requiring explicit synchronization

Message Passing

Communication via explicit messages, simpler to manage but with higher overhead due to kernel involvement

Pipes & FIFOs

Unidirectional communication channels for related (anonymous) or unrelated (named) processes

Sockets

Network-enabled IPC for communication between processes on the same or different machines

Signals

Asynchronous notifications for event handling or process control

Synchronization

Essential for managing concurrent access to shared resources and preventing data inconsistencies

Exam Preparation

Understand Fundamental Purpose

Be clear on the basic operational mechanism of each IPC type

Know the Trade-offs

Be prepared to articulate advantages and disadvantages between different IPC models

Provide Examples

Be ready to give relevant examples for each IPC mechanism

Compare and Contrast

Understand the differences between shared memory and message passing models

 Focus on conceptual understanding rather than memorization

Review Questions

Test your understanding of the key concepts covered in this lecture on Inter-Process Communication (IPC) Mechanisms.

1 IPC Models

What are the main differences between shared memory and message passing?

3 Signals

What role do signals play in IPC?

2 Pipes and FIFOs

How do pipes and FIFOs differ?

4 Synchronization

Why is synchronization important in IPC?