

# File System and Storage Performance

Operating Systems Lecture



Instructor: Tanaya Bowade & Dr Shabih Fatima

# Introduction to File Systems

## What is a File System?

A method and data structure that an operating system uses to control how data is stored and retrieved.

### Key Roles:



#### Data Management

Organizes and manages files and directories on storage devices



#### Abstraction Layer

Provides a logical view of data storage, hiding physical hardware details



#### Data Integrity

Ensures consistency and reliability through journaling and error checking



#### Access Control

Implements security measures like permissions and ownership

## Importance of I/O Performance

Storage and I/O performance directly impact system responsiveness and user experience.

### Why I/O Performance Matters:

- ⌚ Slow I/O operations cause CPU waiting time, reducing system throughput
- ⌚ I/O bottlenecks manifest as slow application loading and unresponsive interfaces
- 💾 Efficient I/O management is critical for modern computing environments
- 📈 Storage technology evolution (HDD to SSD) significantly impacts performance

# File System Concepts

## Core Components



### Files

Basic unit of data storage. A named collection of related information recorded on secondary storage. Can contain text documents, images, or executable programs.



### Directories

Special files that contain lists of other files and subdirectories. Provide hierarchical structure for organizing files, making them easier to locate and manage.



### Metadata

Data about data. Key information includes:



Size



Owner



Permissions



Timestamps



Location



Type

## Attributes & Naming

### File Attributes

#### Read-only

Prevents modification

#### Hidden

Invisible in listings

#### System

Critical OS file

#### Archive

Needs backup

### Naming Conventions

■ **Length Restrictions:** Maximum characters (e.g., 255 in many systems)

■ **Forbidden Characters:** Characters not allowed (e.g., /, \, :, \*, ? in Windows)

■ **Case Sensitivity:** Some systems are case-sensitive (UNIX/Linux)

■ **Extensions:** Suffixes indicating file type (e.g., .txt, .docx)

# File System Structure

## Layered Architecture

File systems are organized in layers, each providing services to the layer above.



### Application Interface

System calls: open(), read(), write(), close()



### Logical File System

File organization, directory management, access control



### File-Organization Module

Logical to physical block mapping, free space management



### Basic File System

I/O scheduling, device-independent I/O



### I/O Control

Device drivers, interrupt handling



### Devices

Disk drives, SSDs, other storage media

## Real-world Implementations

### UNIX File System (UFS)

- ✓ Uses inodes to store metadata separately from data blocks
- ✓ Case-sensitive file names
- ✓ Efficient for sequential access workloads

### NTFS (Windows)

- ✓ Case-insensitive but preserves case
- ✓ Advanced security features and permissions
- ✓ Better support for large volumes

# File Allocation and Access Methods

## File Allocation Methods

### Contiguous Allocation



File occupies consecutive blocks on disk

- +
- Simple implementation
- +
- Fast sequential access
- 
- External fragmentation

### Linked Allocation



File stored as linked or chain list of blocks

- +
- No external fragmentation
- +
- Flexible file size
- 
- Slow direct access

### Indexed Allocation



Index block points to data blocks

- +
- No external fragmentation
- +
- Fast direct access
- 
- Index block overhead

## File Access Methods

### Sequential Access



Data accessed in order, from start to finish

- > Similar to reading a tape or log file
- > Processed one record at a time
- > **Example:** Reading text files line by line

### Direct Access (Random Access)



Access data at any arbitrary position

- > Jump to specific block or record
- > Essential for database applications
- > Requires efficient indexing mechanisms

# Directory Implementation and Management

## Directory Structures

### Single-Level Directory

All files contained in a single directory.

- Directory
- fileA fileB fileC

- + Simplicity
- Naming conflicts

### Two-Level Directory

Separate directory for each user.

- Master Dir
  - User1\_Dir
    - fileA fileB
  - User2\_Dir
    - fileA fileC

- + Resolves naming conflicts
- No user grouping

### Tree-Structured Directory

Hierarchical structure with multiple levels.

- Root Dir1
  - SubDirA
    - fileX
    - fileY

- Dir2
  - SubDirB
    - fileZ

- + Hierarchical organization
- + No naming conflicts

# Storage Devices and Performance Basics


## HDD vs. SSD

Feature	HDD	SSD
Mechanism	Mechanical spinning platters	Flash memory (NAND-based)
Seek Time	Milliseconds (ms)	Microseconds ( $\mu$ s)
Transfer Rate	50-200 MB/s	500-7000 MB/s
Random I/O	Poor	Excellent
Durability	Susceptible to physical shock	More robust, no moving parts

## Key Performance Metrics

- Latency**  
Delay between I/O request and data transfer start. Measured in ms (HDD) or  $\mu$ s (SSD).
- Throughput**  
Amount of data transferred per second (MB/s or GB/s). Critical for sequential operations.
- IOPS**  
Input/output operations per second. Important for workloads with many small random I/Os.

## Performance Comparison



# Storage Performance Monitoring

## Key Metrics

### Latency

Delay between I/O request and data transfer, measured in ms (HDD) or  $\mu$ s (SSD)

### Throughput

Data transferred per unit time, measured in MB/s or GB/s

### IOPS

Input/output operations per second, critical for workloads with many small I/O requests

### Utilization

Percentage of time storage device is busy processing I/O requests

### Queue Depth

Number of I/O requests waiting to be processed

## Monitoring Tools

### Linux Tools

#### iostat - Reports I/O statistics

```
Linux 5.15.0-generic (hostname) 11/03/2025 _x86_64_ (8 CPU)
avg-cpu: %user %nice %system %iowait %steal %idle
0.50 0.00 0.50 0.00 0.00 99.00
Device tps kB_read/s kB_wrtn/s
sda 0.00 0.00 0.00
sdb 0.00 0.00 0.00
```

*tps: Transactions per second*

#### vmstat - Reports system statistics

Includes block I/O statistics (bi, bo) that indicate disk activity

### Windows Tools

#### Performance Monitor (perfmon)

Graphical tool for monitoring storage performance counters

Can track metrics like Disk Queue Length, % Disk Time, and Avg. Disk Response Time

# I/O Performance Analysis

## Analysis Techniques



### Bottleneck Identification

Pinpointing the component in the I/O path that limits overall performance.



### Workload Characterization

- Read-heavy vs. Write-heavy: Proportion of read vs. write operations
- Random vs. Sequential: Access patterns (impact on HDD performance)
- I/O Size: Average size of data blocks being read/written
- Concurrency: Number of simultaneous I/O requests

## Optimization Techniques



### Faster Storage Media

Upgrading from HDD to SSD



### Load Balancing

Distributing I/O across multiple devices



### Caching & Buffering

Implementing effective caching



### RAID Configurations

Combining multiple disks for performance



### Application Improvements

Optimizing code to reduce I/O



### File System Tuning

Adjusting parameters for workloads

# Caching and Buffering

## Role in Performance



### Caching


Stores copies of frequently accessed data in faster, temporary storage (e.g., RAM). Reduces latency for subsequent accesses.



### Buffering

Temporarily holds data during transfer between devices or between a device and an application. Smoothes out differences in data transfer rates.

### Cache Operation:



ⓘ Cache improves read performance by storing frequently accessed data in faster memory

## Modern OS Examples

### Implementation in Modern OS:



#### Linux

Implementing sophisticated caching through the unified buffer cache/page cache, which dynamically manages memory for both file data and metadata.




#### Windows

Uses memory-mapped files and I/O buffering to optimize file system performance and reduce physical disk accesses.

**Key Benefit:** Modern operating systems dynamically allocate and deallocate memory for caching based on system load and application demands, ensuring optimal performance.

# Resource Management Context

File system performance is deeply intertwined with the broader concept of operating system resource management. The OS orchestrates the efficient use of CPU, memory, and I/O devices, which are interconnected and affect each other's performance.



**Key Insight:** Effective resource management requires a holistic approach, as performance issues in one area can impact others.

# Review Questions

Key questions for test preparation



1. Define and explain the main components of a file system.



2. What are the key differences between HDD and SSD performance?



3. How can caching improve I/O performance?



4. Describe one method used to monitor storage performance.



5. How does resource management influence file system efficiency?



These questions cover the key topics. Review your notes and slides for detailed answers.