

Week 8 Journal: File System, Storage Performance, and Security Mechanisms

Course: Operating Systems **Instructors:** Tanaya Bowade & Dr Shabih
Date: November 27, 2025 **Student:** [Your Name]

Table of Contents

1. [File System and Storage Performance](#)
 2. [Introduction to File Systems](#)
 3. [File System Concepts](#)
 4. [File System Structure](#)
 5. [File Allocation and Access Methods](#)
 6. [Directory Implementation and Management](#)
 7. [Storage Devices and Performance Basics](#)
 8. [Storage Performance Monitoring](#)
 9. [I/O Performance Analysis](#)
 10. [Caching and Buffering](#)
 11. [Resource Management Context](#)
 12. [Security Mechanisms and Access Control](#)
 13. [Learning Objectives](#)
 14. [Mandatory Access Control Implementation](#)
 15. [Intrusion Detection with fail2ban](#)
 16. [Security Verification and Monitoring Scripts](#)
 17. [Review Questions](#)
 18. [Reflection](#)
-

File System and Storage Performance

Introduction to File Systems

What is a File System?

A **file system** is a method and data structure that an operating system uses to control how data is stored and retrieved. It provides essential functionality for managing data on storage devices.

Key Roles of a File System:

1. Data Management

2. Organizes and manages files and directories on storage devices
3. Provides hierarchical structure for efficient data organization

4. Abstraction Layer

5. Provides a logical view of data storage
6. Hides physical hardware details from users and applications
7. Makes storage access uniform regardless of underlying hardware

8. Data Integrity

9. Ensures consistency and reliability through journaling
10. Implements error checking mechanisms

11. Protects against data corruption

12. Access Control

13. Implements security measures
14. Manages permissions and ownership
15. Controls who can read, write, or execute files

Importance of I/O Performance

Storage and I/O performance directly impact system responsiveness and user experience.

Why I/O Performance Matters:

- **CPU Waiting Time:** Slow I/O operations cause CPU waiting time, reducing system throughput
 - **User Experience:** I/O bottlenecks manifest as slow application loading and unresponsive interfaces
 - **Critical Operations:** Efficient I/O management is critical for modern computing environments
 - **Technology Evolution:** Storage technology evolution (HDD to SSD) significantly impacts performance
-

File System Concepts

Core Components

1. Files - Basic unit of data storage - A named collection of related information recorded on secondary storage - Can contain text documents, images, executable programs, or any type of data - Identified by unique names within their directories

2. Directories - Special files that contain lists of other files and subdirectories - Provide hierarchical structure for organizing files - Make files easier to locate and manage - Can be nested to create complex directory trees

3. Metadata - Data about data - information describing file characteristics - Key metadata includes: - **Size:** Amount of space the file occupies - **Owner:** User who owns the file - **Permissions:** Access rights for different users - **Timestamps:** Creation, modification, and access times - **Location:** Physical location on storage device - **Type:** File format or content type

Attributes & Naming

File Attributes:

1. **Read-only**: Prevents modification of the file
2. **Hidden**: Makes file invisible in standard directory listings
3. **System**: Marks file as critical OS file
4. **Archive**: Indicates file needs backup

Naming Conventions:

1. **Length Restrictions**: Maximum characters (e.g., 255 in many systems)
 2. **Forbidden Characters**: Characters not allowed (e.g., /, \, :, *, ? in Windows)
 3. **Case Sensitivity**: Some systems are case-sensitive (UNIX/Linux), others are not (Windows)
 4. **Extensions**: Suffixes indicating file type (e.g., .txt, .docx, .pdf)
-

File System Structure

Layered Architecture

File systems are organized in layers, each providing services to the layer above:

1. **Application Interface**
2. System calls: open(), read(), write(), close()
3. Provides programming interface for file operations
4. **Logical File System**
5. File organization and directory management
6. Access control and security
7. Manages file metadata
8. **File-Organization Module**
9. Logical to physical block mapping
10. Free space management
11. File allocation strategies

12. Basic File System

13. I/O scheduling
14. Device-independent I/O operations

15. Buffer management

16. I/O Control

17. Device drivers
18. Interrupt handling
19. Direct hardware communication

20. Devices

21. Physical storage: Disk drives, SSDs, other storage media
22. Hardware controllers

Real-world Implementations

UNIX File System (UFS): - Uses inodes to store metadata separately from data blocks - Case-sensitive file names - Efficient for sequential access workloads - Simple and robust design

NTFS (Windows): - Case-insensitive but preserves case - Advanced security features and permissions - Better support for large volumes - Journaling capability for reliability - Support for file compression and encryption

File Allocation and Access Methods

File Allocation Methods

1. Contiguous Allocation - File occupies consecutive blocks on disk -
Advantages: - Simple implementation - Fast sequential access - Minimal seek time - **Disadvantages:** - External fragmentation - Difficult to grow files - Need to know file size in advance

2. Linked Allocation - File stored as linked or chain list of blocks - Each block contains pointer to next block - **Advantages:** - No external fragmentation - Flexible file size - Easy to grow files - **Disadvantages:** - Slow direct access - Pointer overhead in each block - Reliability issues if pointer corrupted

3. Indexed Allocation - Index block points to data blocks - All pointers to file blocks kept in index block - **Advantages:** - No external fragmentation - Fast direct access - Supports random access efficiently - **Disadvantages:** - Index block overhead - Wasted space for small files

File Access Methods

1. Sequential Access - Data accessed in order, from start to finish - Similar to reading a tape or log file - Processed one record at a time - **Example:** Reading text files line by line, processing logs

2. Direct Access (Random Access) - Access data at any arbitrary position - Jump to specific block or record - Essential for database applications - Requires efficient indexing mechanisms - **Example:** Database queries, multimedia file access

Directory Implementation and Management

Directory Structures

1. Single-Level Directory - All files contained in a single directory - **Advantages:** - Simplicity in implementation - Easy to understand - **Disadvantages:** - Naming conflicts between users - No organization capability - Not scalable

2. Two-Level Directory - Separate directory for each user - Master directory contains user directories - **Advantages:** - Resolves naming conflicts between users - Basic isolation between users - **Disadvantages:** - No grouping capability for single user - Limited organization

3. Tree-Structured Directory - Hierarchical structure with multiple levels - Directories can contain files and subdirectories - **Advantages:** - Hierarchical organization - No naming conflicts - Flexible structure -

Efficient searching with paths - **Disadvantages:** - More complex implementation - Potential for deep nesting

Storage Devices and Performance Basics

HDD vs. SSD Comparison

Feature	HDD	SSD
Mechanism	Mechanical spinning platters	Flash memory (NAND-based)
Seek Time	Milliseconds (ms)	Microseconds ($\frac{1}{\text{ms}}$)
Transfer Rate	50-200 MB/s	500-7000 MB/s
Random I/O	Poor	Excellent
Durability	Susceptible to physical shock	More robust, no moving parts
Lifespan	Mechanical wear	Write cycle limitations
Cost	Lower cost per GB	Higher cost per GB
Noise	Audible operation	Silent operation

Key Performance Metrics

- 1. Latency** - Delay between I/O request and data transfer start - Measured in ms (HDD) or $\frac{1}{\text{ms}}$ (SSD) - Critical for responsiveness - Affects user-perceived performance
- 2. Throughput** - Amount of data transferred per second - Measured in MB/s or GB/s - Critical for sequential operations - Important for large file transfers
- 3. IOPS (Input/Output Operations Per Second)** - Number of I/O operations completed per second - Important for workloads with many small random I/Os - Database performance heavily dependent on IOPS - SSDs excel in IOPS compared to HDDs

Performance Comparison: - **Seek Time:** SSD ~100x faster than HDD - **Transfer Rate:** SSD 10-35x faster than HDD - **Random I/O:** SSD significantly outperforms HDD

Storage Performance Monitoring

Key Metrics

- 1. Latency** - Delay between I/O request and data transfer - Measured in ms (HDD) or μs (SSD) - Lower latency = better responsiveness
- 2. Throughput** - Data transferred per unit time - Measured in MB/s or GB/s - Higher throughput = better sequential performance
- 3. IOPS** - Input/output operations per second - Critical for workloads with many small I/O requests - Database and transactional workloads
- 4. Utilization** - Percentage of time storage device is busy processing I/O requests - High utilization may indicate bottleneck
- 5. Queue Depth** - Number of I/O requests waiting to be processed - High queue depth may indicate performance issues

Monitoring Tools

Linux Tools:

1. **iostat** - Reports I/O statistics `bash iostat -x 1 5`
2. Shows device utilization
3. Reports tps (transactions per second)
4. Displays read/write rates
5. **vmstat** - Reports system statistics `bash vmstat 1 10`
6. Includes block I/O statistics (bi, bo)
7. Shows memory and CPU usage
8. Indicates disk activity

Windows Tools:

- 1. Performance Monitor (perfmon)**
 2. Graphical tool for monitoring storage performance counters
 3. Can track metrics like:
 - Disk Queue Length
 - % Disk Time
 - Avg. Disk Response Time
 4. Allows creating custom monitoring views
-

I/O Performance Analysis

Analysis Techniques

1. Bottleneck Identification - Pinpointing the component in the I/O path that limits overall performance - Analyzing CPU wait time for I/O - Examining queue depths - Identifying saturation points

2. Workload Characterization

Understanding the nature of I/O operations:

- **Read-heavy vs. Write-heavy:** Proportion of read vs. write operations
- **Random vs. Sequential:** Access patterns (significant impact on HDD performance)
- **I/O Size:** Average size of data blocks being read/written
- **Concurrency:** Number of simultaneous I/O requests

Optimization Techniques

1. Faster Storage Media - Upgrading from HDD to SSD - Using NVMe drives for even better performance - Tiered storage strategies

2. RAID Configurations - Combining multiple disks for performance - RAID 0 for performance (striping) - RAID 10 for performance and redundancy

3. Load Balancing - Distributing I/O across multiple devices -
Preventing single device saturation - Using multiple I/O paths

4. Application Improvements - Optimizing code to reduce I/O
operations - Batching small operations - Using asynchronous I/O

5. Caching & Buffering - Implementing effective caching strategies -
Using write-back caching - Optimizing buffer sizes

6. File System Tuning - Adjusting parameters for specific workloads -
Choosing appropriate block sizes - Optimizing journal settings

Caching and Buffering

Role in Performance

Caching: - Stores copies of frequently accessed data in faster, temporary storage (e.g., RAM) - Reduces latency for subsequent accesses -
Improves read performance significantly - Based on temporal and spatial locality principles

Cache Operation Flow: 1. **Cache Hit:** Requested data found in cache (fast access)
2. **Cache Miss:** Data not in cache, must fetch from slower storage
3. Cache improves performance by storing frequently accessed data in faster memory

Buffering: - Temporarily holds data during transfer between devices -
Smoothes out differences in data transfer rates - Reduces number of physical I/O operations - Allows asynchronous operation

Modern OS Examples

Implementation in Modern OS:

Linux: - Implements sophisticated caching through the unified buffer cache/page cache - Dynamically manages memory for both file data and metadata - Uses algorithms like LRU (Least Recently Used) for cache management - Automatically adjusts cache size based on system load

Windows: - Uses memory-mapped files for efficient file access -
Implements I/O buffering to optimize file system performance - Reduces physical disk accesses through intelligent caching - Lazy write mechanism for improved write performance

Key Benefit: Modern operating systems dynamically allocate and deallocate memory for caching based on system load and application demands, ensuring optimal performance without manual configuration.

Resource Management Context

File system performance is deeply intertwined with the broader concept of operating system resource management. The OS orchestrates the efficient use of CPU, memory, and I/O devices, which are interconnected and affect each other's performance.

CPU Management Impact on File System

- **High CPU Utilization:** CPU busy with other processes delays file system request processing
- **Insufficient CPU Cycles:** Slow metadata updates and file operations
- **Process Scheduling:** Affects I/O request handling priority

Memory Management Impact on File System

- **Excessive Page Swapping:** Increases I/O operations and slows system performance
- **Limited Memory for Caches:** Reduces cache hit rate, forcing more disk accesses
- **Memory Pressure:** Reduces buffer cache effectiveness

I/O Management Impact on File System

- **Device Drivers:** Quality of drivers impacts I/O request processing
- **Storage Controllers:** Hardware capabilities affect performance
- **Queue Depth:** Affects number of pending I/O requests

- **Scheduling Algorithms:** Determines I/O request ordering and fairness

Key Insight: Effective resource management requires a holistic approach, as performance issues in one area can impact others. Optimizing file system performance requires considering the entire system context.

Security Mechanisms and Access Control

Learning Objectives

1. Implement mandatory access control systems (AppArmor/SELinux)
2. Configure intrusion detection tools (fail2ban)
3. Develop security verification scripts
4. Create remote monitoring capabilities

Pre-Lab Preparation

Required Knowledge: - Mandatory Access Control Concepts - Introduction to Intrusion Detection Systems - Bash Scripting for System Administration - Review: Operating System Modes

Mandatory Access Control Implementation

Understanding Access Control Systems

Discretionary Access Control (DAC): - Traditional file permissions (user, group, others) - Owner has discretion over permissions - Vulnerable to privilege escalation

Mandatory Access Control (MAC): - System-wide security policy - Users cannot override security policies - More restrictive and secure - Used in high-security environments

Access Control Systems

Ubuntu/Debian: - Typically uses **AppArmor** - Profile-based security

Red Hat/CentOS: - Typically uses **SELinux** - Label-based security

Task 1.1: Identifying Your Access Control System

Check which system is active:

```
```bash
```

## **For AppArmor (Ubuntu/Debian)**

```
sudo aa-status
```

## **For SELinux (Red Hat/CentOS)**

```
sestatus ```
```

### **Task 1.2: Working with AppArmor Profiles**

**1. List all AppArmor profiles and their modes:** bash sudo aa-status

**2. Examine a specific profile:** bash sudo cat  
/etc/apparmor.d/usr.sbin.tcpdump

**3. View profiles in different modes:** bash sudo aa-status --  
profiled sudo aa-status --enforced sudo aa-status --  
complaining

### **Understanding Profile Modes:**

- **Enforce Mode:**
- Violations are blocked and logged
- Provides active protection
- Most secure setting
- **Complain Mode:**

- Violations are logged but allowed
- Used for testing and profiling
- Helps develop new profiles

### **Task 1.3: Creating an AppArmor Status Report Script**

**Purpose:** Automated reporting of AppArmor configuration and status

**Script:** apparmor-report.sh

```bash

!/bin/bash

AppArmor Status Report Script

**Reports on all AppArmor profiles
and their status**

```
echo "===== echo
"AppArmor Status Report" echo
===== echo
"Generated: $(date)" echo "Hostname: $(hostname)" echo ""
```

Check if AppArmor is installed

```
if ! command -v aa-status &> /dev/null; then echo "ERROR: AppArmor
is not installed" exit 1 fi
```

```
echo "==== Profile Summary ==="
```

Count total profiles loaded

```
total_profiles=$(sudo aa-status --profiled | wc -l) echo "Total profiles loaded: $total_profiles" echo ""

echo "==== Enforced Profiles ====" sudo aa-status --enforced
enforced_count=$(sudo aa-status --enforced | wc -l) echo "Count: $enforced_count" echo ""

echo "==== Complain Mode Profiles ====" sudo aa-status --complaining
complain_count=$(sudo aa-status --complaining | wc -l) echo "Count: $complain_count" echo ""

echo "===== Report Complete =====" echo
"===== " ``
```

Make executable and run: bash chmod +x apparmor-report.sh
./apparmor-report.sh

For SELinux Systems

Task 1.2: Understanding SELinux Status

```bash

## Check SELinux status

sestatus getenforce

## View security contexts

ls -Z /etc/ssh/sshd\_config ps -eZ | grep sshd

## Check for recent denials

sudo ausearch -m avc -ts recent ```

**SELinux Security Contexts:** - Every file, process, and resource has a security label - Format: user:role:type:level - Policies define what operations are allowed

---

## Intrusion Detection with fail2ban

### What is fail2ban?

- Intrusion prevention software
- Monitors log files for malicious activity
- Automatically bans IP addresses showing malicious signs
- Protects against brute-force attacks

### Task 2.1: Installing and Configuring fail2ban

**1. Install fail2ban:** bash sudo apt update sudo apt install fail2ban

**2. Check service status:** bash sudo systemctl status fail2ban

**3. Create local configuration:** bash sudo cp /etc/fail2ban/jail.conf /etc/fail2ban/jail.local

**Important:** Never edit jail.conf directly; always use jail.local

**4. Configure SSH protection:** bash sudo nano /etc/fail2ban/jail.local

**Configuration in [sshd] section:** ini [sshd] enabled = true port = 22 filter = sshd logpath = /var/log/auth.log maxretry = 3 bantime = 600 findtime = 600

### Parameters Explained:

- **maxretry:** Number of failed attempts before ban (e.g., 3)
- **bantime:** How long to ban in seconds (e.g., 600 = 10 minutes)
- **findtime:** Time window for counting failures in seconds (e.g., 600 = 10 minutes)

**Logic:** If 3 failed attempts occur within 10 minutes, ban the IP for 10 minutes.

**5. Restart and enable fail2ban:** bash sudo systemctl restart fail2ban sudo systemctl enable fail2ban

### Task 2.2: Monitoring fail2ban Activity

**1. Check active jails:** bash sudo fail2ban-client status sudo fail2ban-client status sshd

**2. View fail2ban log:** bash sudo tail -30 /var/log/fail2ban.log

**3. View banned IPs:** bash sudo fail2ban-client banned

**4. Manually unban an IP if needed:** bash sudo fail2ban-client set sshd unbanip [IP\_ADDRESS]

**5. Examine authentication logs:** bash sudo grep "Failed password" /var/log/auth.log | tail -10

### Task 2.3: Configuring Automatic Security Updates

**1. Install unattended-upgrades:** bash sudo apt install unattended-upgrades

**2. Enable automatic security updates:** bash sudo dpkg-reconfigure -plow unattended-upgrades Select "Yes" when prompted

**3. Verify service status:** bash sudo systemctl status unattended-upgrades

**4. Check configuration:** bash sudo cat /etc/apt/apt.conf.d/50unattended-upgrades

**5. View update logs:** bash sudo cat /var/log/unattended-upgrades/unattended-upgrades.log

**Security vs Stability Trade-off:** - **Pro:** Automatic security patches protect against vulnerabilities - **Con:** Updates might introduce

compatibility issues - **Recommendation:** Enable for security updates, test major updates manually

---

## Security Verification and Monitoring Scripts

### Task 3.1: Creating the Security Baseline Verification Script

**Purpose:** Comprehensive verification of all security configurations

**Script:** `security-baseline.sh`

This script verifies: - SSH security configuration - Firewall status and rules - Intrusion detection (fail2ban) - Mandatory access control (AppArmor/SELinux) - Administrative users - Automatic security updates - System resources

**Key Features:** - Color-coded output for easy reading - Comprehensive security checks - Identifies security warnings - Reports on all critical security components

**Script sections:** 1. SSH configuration verification 2. Firewall status check 3. fail2ban monitoring 4. MAC system verification 5. User privilege audit 6. Automatic update status 7. System resource monitoring

### Task 3.2: Creating the Remote Monitoring Script

**Purpose:** Monitor server remotely from workstation via SSH

**Script:** `monitor-server.sh`

**Capabilities:** - System information collection - Uptime and load monitoring - CPU usage analysis - Memory utilization tracking - Disk usage and I/O statistics - Network connection monitoring - Failed login attempt tracking - fail2ban status reporting

**Key Features:** - Runs from workstation - Connects via SSH - Logs all output to file - Error handling for connection failures - Comprehensive metric collection

**Metrics Collected:** 1. System information (OS, kernel version) 2. Uptime and load average 3. Top CPU-consuming processes 4. Memory usage (free -h) 5. Disk usage (df -h) 6. Disk I/O statistics (iostat) 7. Network connection summary 8. Active network connections 9. Recent failed login attempts 10. fail2ban jail status

---

## Review Questions

### File System and Storage Performance

#### 1. Define and explain the main components of a file system.

A file system consists of three main components:

- **Files:** Basic unit of data storage; a named collection of related information recorded on secondary storage. Files can contain text documents, images, or executable programs.
- **Directories:** Special files that contain lists of other files and subdirectories. They provide hierarchical structure for organizing files, making them easier to locate and manage.
- **Metadata:** Data about data. Key information includes size, owner, permissions, timestamps, location, and type. Metadata enables the file system to manage and track files efficiently.

#### 2. What are the key differences between HDD and SSD performance?

Key differences include:

- **Mechanism:** HDDs use mechanical spinning platters; SSDs use flash memory (NAND-based)
- **Seek Time:** HDDs have millisecond seek times; SSDs have microsecond seek times (~100x faster)
- **Transfer Rate:** HDDs: 50-200 MB/s; SSDs: 500-7000 MB/s
- **Random I/O:** HDDs have poor random I/O performance; SSDs excel at random I/O

- **Durability:** HDDs are susceptible to physical shock; SSDs are more robust with no moving parts

### **3. How can caching improve I/O performance?**

Caching improves I/O performance by:

- Storing copies of frequently accessed data in faster temporary storage (RAM)
- Reducing latency for subsequent accesses to the same data
- Eliminating need for slow disk access on cache hits
- Leveraging temporal and spatial locality principles
- Modern OS dynamically manages cache size based on system load

When data is requested, the system first checks the cache. If found (cache hit), data is returned quickly from RAM. If not found (cache miss), data is fetched from slower storage and added to cache for future access.

### **4. Describe one method used to monitor storage performance.**

**iostat (Linux tool):**

iostat reports I/O statistics for storage devices. It provides: - Device utilization percentage - Transactions per second (tps) - Read and write rates (kB\_read/s, kB\_wrtn/s) - Average wait times and service times

Usage: `iostat -x 1 5` displays extended statistics, updated every 1 second, 5 times.

This tool helps identify I/O bottlenecks by showing which devices are heavily utilized and how much data is being transferred. High utilization or long wait times indicate potential performance issues.

### **5. How does resource management influence file system efficiency?**

Resource management influences file system efficiency through interconnected components:

- **CPU Management:** High CPU utilization delays file system request processing; insufficient CPU cycles slow metadata updates

- **Memory Management:** Excessive page swapping increases I/O operations; limited memory for caches reduces hit rate, forcing more disk accesses
- **I/O Management:** Device drivers and storage controllers impact I/O request processing; queue depth and scheduling algorithms affect storage performance

Effective resource management requires a holistic approach because performance issues in one area (e.g., memory pressure) directly impact file system performance by reducing cache effectiveness and increasing disk I/O.

## **Security Mechanisms and Access Control**

### **6. What is the difference between Mandatory Access Control (MAC) and Discretionary Access Control (DAC)?**

**Discretionary Access Control (DAC):** - Traditional Unix/Linux file permissions (user, group, others) - File owners have discretion over who can access their files - Users can modify permissions on files they own - More flexible but less secure - Vulnerable to privilege escalation attacks

**Mandatory Access Control (MAC):** - System-wide security policy enforced by the OS - Users cannot override security policies, even on files they own - Based on security labels and policies - More restrictive and secure - Implemented through AppArmor (Ubuntu) or SELinux (Red Hat) - Used in high-security environments

### **7. Explain how fail2ban protects against brute-force attacks.**

fail2ban protects against brute-force attacks by:

1. **Monitoring:** Continuously monitors log files (e.g., /var/log/auth.log) for failed login attempts
2. **Detection:** Identifies patterns of malicious behavior (e.g., repeated failed passwords)
3. **Action:** Automatically bans IP addresses that exceed the threshold for failed attempts
4. **Temporary Ban:** Bans are temporary (configured by bantime), allowing legitimate users who made mistakes to try again later

**Configuration parameters:** - **maxretry**: Number of failures before ban (e.g., 3) - **findtime**: Time window for counting failures (e.g., 600 seconds) - **bantime**: Duration of ban (e.g., 600 seconds)

If an IP has 3 failed attempts within 10 minutes, it gets banned for 10 minutes, making brute-force attacks impractical.

## **8. Why are automatic security updates important, and what is the trade-off?**

**Importance:** - Protect against newly discovered vulnerabilities - Reduce window of exposure to attacks - Ensure system has latest security patches - Reduce administrative burden

**Trade-offs:** - **Security vs. Stability**: Updates might introduce bugs or compatibility issues - **Automation vs. Control**: Less control over what gets updated and when - **Testing**: No opportunity to test updates before deployment

**Best Practice:** Enable automatic security updates while manually testing major system updates in a test environment first.

---

## **Reflection**

### **Key Learnings**

This week covered two critical aspects of operating systems:

**1. File System and Storage Performance:** - Understanding file system architecture and layers - Recognizing the importance of I/O performance optimization - Learning about different storage technologies (HDD vs SSD) - Mastering performance monitoring and analysis techniques - Implementing caching and buffering strategies

**2. Security Mechanisms:** - Implementing mandatory access control (AppArmor/SELinux) - Configuring intrusion detection with fail2ban - Creating automated security verification scripts - Developing remote monitoring capabilities - Understanding the defense-in-depth security model

## **Challenges and Solutions**

**Challenge 1: Understanding MAC vs DAC** - MAC systems add complexity but significantly enhance security - Profiles/policies require careful configuration to avoid breaking applications

**Challenge 2: Balancing Security and Usability** - Strict security measures can impact user experience - Finding the right balance between security and functionality is crucial

**Challenge 3: Performance Monitoring** - Multiple metrics (latency, throughput, IOPS) must be considered together - Context matters - what's good for one workload may be poor for another

## **Practical Applications**

**1. Performance Optimization:** - Use SSDs for databases and high-I/O applications - Implement caching for frequently accessed data - Monitor I/O metrics to identify bottlenecks - Choose appropriate file allocation methods based on workload

**2. Security Hardening:** - Layer multiple security mechanisms (firewall + fail2ban + MAC) - Automate security verification with scripts - Implement remote monitoring for proactive security - Keep systems updated with automatic security patches

## **Connection to Learning Outcomes**

**LO3 - System Administration:** - Configuring and managing security systems - Implementing monitoring and verification scripts - Understanding system performance metrics

**LO4 - Security and Performance:** - Balancing security requirements with performance needs - Implementing defense-in-depth security strategies - Optimizing I/O performance through caching and proper storage selection

## **Future Considerations**

- 1. Emerging Technologies:**
2. NVMe storage for even better performance
3. Container security with AppArmor/SELinux
4. Advanced monitoring with AI/ML-based anomaly detection

### **5. Best Practices:**

6. Regular security audits using baseline verification scripts
7. Continuous performance monitoring
8. Proactive security updates
9. Documentation of all security configurations

### **10. Areas for Further Study:**

11. Advanced file system features (snapshots, deduplication)
  12. Performance tuning for specific workloads
  13. Security information and event management (SIEM)
  14. Incident response procedures
- 

## **Conclusion**

Week 8 provided comprehensive coverage of file system concepts, storage performance, and security mechanisms. The integration of theoretical knowledge with practical implementation through scripts and monitoring tools reinforces the importance of both understanding underlying concepts and developing practical skills.

The key takeaway is that effective system administration requires: - Deep understanding of how file systems work - Ability to measure and optimize performance - Implementation of layered security controls - Automation of monitoring and verification tasks - Continuous improvement through monitoring and analysis

These skills are essential for managing modern computing systems that must balance performance, security, and reliability requirements.

---

## **End of Week 8 Journal**

# Table of Contents

|                                              |    |
|----------------------------------------------|----|
| File System and Storage Performance          | 2  |
| Introduction to File Systems                 | 2  |
| File System Concepts                         | 3  |
| File System Structure                        | 4  |
| File Allocation and Access Methods           | 5  |
| Directory Implementation and Management      | 6  |
| Storage Devices and Performance Basics       | 7  |
| Storage Performance Monitoring               | 8  |
| I/O Performance Analysis                     | 9  |
| Caching and Buffering                        | 10 |
| Resource Management Context                  | 11 |
| Security Mechanisms and Access Control       | 12 |
| Learning Objectives                          | 12 |
| Mandatory Access Control Implementation      | 12 |
| Intrusion Detection with fail2ban            | 16 |
| Security Verification and Monitoring Scripts | 18 |
| Review Questions                             | 19 |
| Reflection                                   | 22 |